

Debugging

Computer systems are **really REALLY complicated**
because it consisting of many moving parts
working together.

If something goes **wrong**, how do we know what
and where to look for the cause?

Debugging is the process of finding and resolving bugs within computer programs, software or systems.

- Wikipedia

Debugging is similar being your local GP doctor.

You are trying to ‘guess’ the problem using your own methods, tools, intuition, knowledge and skills before providing ‘treatment’.
(or referring to a ‘specialist’)

Types of errors

Compile-time Errors!

Compile-time errors are errors that occur when you build a program with a compiler.

Compiler-time errors are great because they help catch invalid lines before the program can be run or shipped

Not all languages have compilers!!

Run-time errors

Run-time errors are errors that happen while the program is running, usually in a form of a crash.

Run-time errors can be forced by:

- **Developer** (in some form of exceptions)
 - **OS** (e.g. segmentation fault)

Crashing is always bad, of course, but be aware that sometimes it's better to crash than to let the program continue.

Crashing is always bad, of course, but be aware that sometimes it's better to crash than to let the program continue.

(i.e. crash bad but not crash maybe worse)

Logical errors

Logical errors are when the program does something that you **do not** expect it to do.

Logical errors are the hardest to debug
because it depends on what the error is.

Usually, we will need to rely on some tool to
help us find the bug.

User error prevention and catching

Some kind of error code

The simplest kind of run-time error checking is to simply output **something** from a function that **indicates the type of error**, so that the person using it can handle it accordingly.


```

2// Some APIs do something like this...
3// I'm looking at you WinAPI and OpenGL and STD
4
5// Hidden from users
6typedef enum GLError {
7    GL_NO_ERROR,
8    GL_INVALID_ENUM,
9    GL_INVALID_VALUE,
10    // etc...
11} GLError;
12
13static GLError error = GL_NO_ERROR;
14

```

```

14// Visible to users
15static int FuncThatReturnTrueFalse(...) {
16    if (!something) {
17        error = GL_INVALID_VALUE;
18        return 0;
19    }
20    if (!something_else(...)) {
21        error = GL_INVALID_ENUM;
22        return 0;
23    }
24    return 1;
25}
26
27static GLError GetLastError() {
28    return error;
29}
30
31int main() {
32    if (!FuncThatReturnTrueFalse(...)) {
33        GLError err = GetLastError();
34        switch(err) {
35            //...
36        }
37    }
38
39} int main() {

```

Asserts

Asserts will crash the program when a line is encountered.

Some assert implementations return the **line number and file** in which the assert occurred, which is useful!


```
momo@DESKTOP-N6DP5P1:/mnt/d/work/sandbox/c$ gcc main.c -o main.exe && ./main.exe
accessing arr[0]
accessing arr[1]
accessing arr[2]
accessing arr[3]
accessing arr[4]
accessing arr[5]
main.exe: main.c:25: GetArrayInt: Assertion `index < arr->count' failed.
Aborted (core dumped)
momo@DESKTOP-N6DP5P1:/mnt/d/work/sandbox/c$ |
```

Typically used in **development** builds to catch statements that should 'never' happen.

Usually, it's **disabled** when shipping a **production** build.

Note that there is some overhead when using asserts. It's like spamming 'if' conditions everywhere.

But sometimes,
it's better to be safe than slow.

Keyword: sometimes.

Unit testing

Writing code that test code.

But what tests the code that test the code?

**But what tests the code that tests the code that
tests the code that tests the code that tests the
code that tests the code that tests the code
that tests the code that tests the code that tests
the code that tests the code that tests the code
that tests the code that tests the code that tests
the code that tests the code that tests the code
that tests the code that tests the code that tests
the code that tests the code that tests the code
that tests the code that tests the code?**

Can all features be tested with code?

If all tests **pass**, does that mean we are **safe**?

In the end, it **depends**.

Always weigh the pros and cons.

Always consider the feasibility.

Always remember that test code is still code
to be written and maintained.

Archiving

Archiving is to write data out to a file, so that when an error occurs, we can trace what happened at when.

We typically archive:

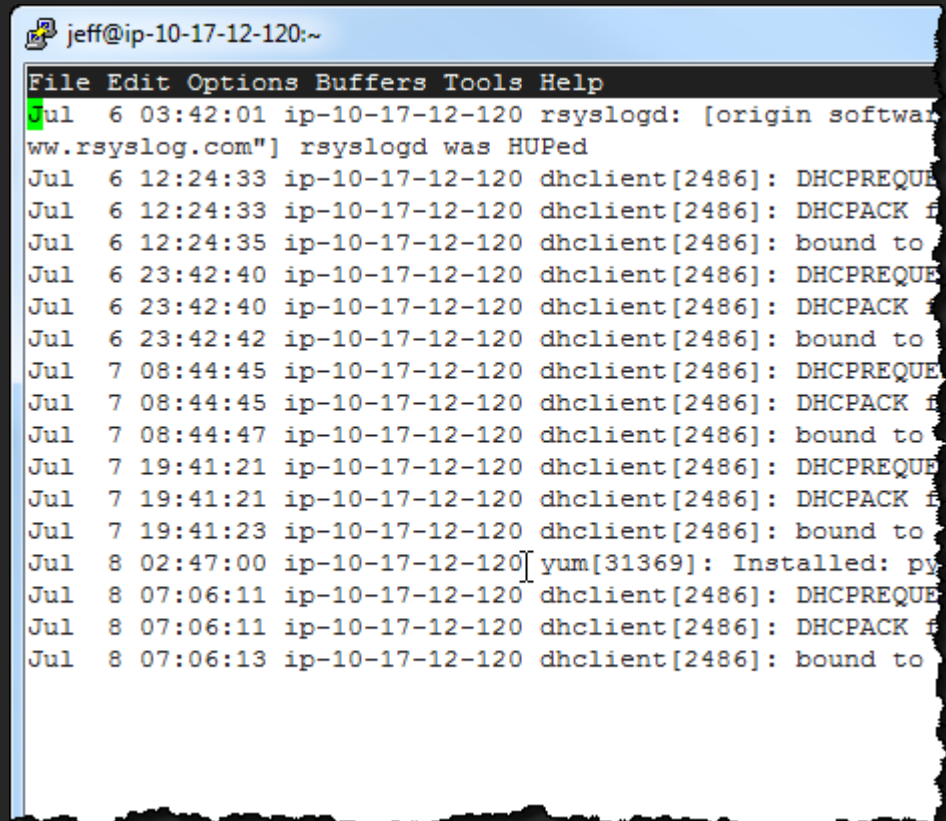
- **Logs**, text files written as program runs
- **Dumps**, snapshots of state (e.g. memory)

Logging

Logging is simply writing text to a file
'just in case'

```
29int main() {  
30  
31    if (...) {  
32        printf("There is a problem here!");  
33    }  
34  
35    if (...) {  
36        printf("And here!!");  
37    }  
38  
39    if (...) {  
40        printf("And also here!");  
41    }  
42}  
43
```

Logging into human readable text is the simplest yet most flexible form of debugging.

A screenshot of a terminal window with a blue title bar. The title bar contains a small icon and the text 'jeff@ip-10-17-12-120:~'. The terminal has a menu bar with 'File Edit Options Buffers Tools Help'. The main area displays a series of log entries. The first entry is highlighted with a green cursor. The logs show various system events including rsyslogd being HUPed, DHCP transactions between ip-10-17-12-120 and a client, and yum installing a package.

```
jeff@ip-10-17-12-120:~  
File Edit Options Buffers Tools Help  
Jul  6 03:42:01 ip-10-17-12-120 rsyslogd: [origin software  
ww.rsyslog.com"] rsyslogd was HUPed  
Jul  6 12:24:33 ip-10-17-12-120 dhclient[2486]: DHCPREQUEST  
Jul  6 12:24:33 ip-10-17-12-120 dhclient[2486]: DHCPACK f  
Jul  6 12:24:35 ip-10-17-12-120 dhclient[2486]: bound to  
Jul  6 23:42:40 ip-10-17-12-120 dhclient[2486]: DHCPREQUEST  
Jul  6 23:42:40 ip-10-17-12-120 dhclient[2486]: DHCPACK f  
Jul  6 23:42:42 ip-10-17-12-120 dhclient[2486]: bound to  
Jul  7 08:44:45 ip-10-17-12-120 dhclient[2486]: DHCPREQUEST  
Jul  7 08:44:45 ip-10-17-12-120 dhclient[2486]: DHCPACK f  
Jul  7 08:44:47 ip-10-17-12-120 dhclient[2486]: bound to  
Jul  7 19:41:21 ip-10-17-12-120 dhclient[2486]: DHCPREQUEST  
Jul  7 19:41:21 ip-10-17-12-120 dhclient[2486]: DHCPACK f  
Jul  7 19:41:23 ip-10-17-12-120 dhclient[2486]: bound to  
Jul  8 02:47:00 ip-10-17-12-120 yum[31369]: Installed: py  
Jul  8 07:06:11 ip-10-17-12-120 dhclient[2486]: DHCPREQUEST  
Jul  8 07:06:11 ip-10-17-12-120 dhclient[2486]: DHCPACK f  
Jul  8 07:06:13 ip-10-17-12-120 dhclient[2486]: bound to
```

You can log pretty much **anything and everything**! It is not just for debugging.

Just remember to **adhere to a format** so that it's easier to filter the texts with the tools you use.

```
2015-10-17 15:45:11,258 INFO [main] org.apache.hadoop.metrics2.impl.MetricsConfig: loaded properties from hadoop-metrics2.properties
2015-10-17 15:45:11,399 INFO [main] org.apache.hadoop.metrics2.impl.MetricsSystemImpl: Scheduled snapshot period at 10 second(s).
2015-10-17 15:45:11,399 INFO [main] org.apache.hadoop.metrics2.impl.MetricsSystemImpl: MapTask metrics system started
2015-10-17 15:45:11,430 INFO [main] org.apache.hadoop.mapred.YarnChild: Executing with tokens:
2015-10-17 15:45:11,430 INFO [main] org.apache.hadoop.mapred.YarnChild: Kind: mapreduce.job, Service: job_1445062781478_0015, Ident: (org.apache$
2015-10-17 15:45:11,602 INFO [main] org.apache.hadoop.mapred.YarnChild: Sleeping for 0ms before retrying again. Got null now.
2015-10-17 15:45:12,196 INFO [main] org.apache.hadoop.mapred.YarnChild: mapreduce.cluster.local.dir for child: /tmp/hadoop-msrabi/nm-local-dir/u$
2015-10-17 15:45:12,711 INFO [main] org.apache.hadoop.conf.Configuration.deprecation: session.id is deprecated. Instead, use dfs.metrics.session$
2015-10-17 15:45:13,602 INFO [main] org.apache.hadoop.yarn.util.ProcfsBasedProcessTree: ProcfsBasedProcessTree currently is supported only on Li$
2015-10-17 15:45:13,618 INFO [main] org.apache.hadoop.mapred.Task: Using ResourceCalculatorProcessTree : org.apache.hadoop.yarn.util.WindowsBas$
2015-10-17 15:45:14,008 INFO [main] org.apache.hadoop.mapred.MapTask: Processing split: hdfs://msra-sa-41:9000/pageinput2.txt:402653184+134217728
2015-10-17 15:45:14,102 INFO [main] org.apache.hadoop.mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
2015-10-17 15:45:14,102 INFO [main] org.apache.hadoop.mapred.MapTask: mapreduce.task.io.sort.mb: 100
2015-10-17 15:45:14,102 INFO [main] org.apache.hadoop.mapred.MapTask: soft limit at 83886080
2015-10-17 15:45:14,102 INFO [main] org.apache.hadoop.mapred.MapTask: bufstart = 0; bufvoid = 104857600
2015-10-17 15:45:14,102 INFO [main] org.apache.hadoop.mapred.MapTask: kvstart = 26214396; length = 6553600
2015-10-17 15:45:14,118 INFO [main] org.apache.hadoop.mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuf$
2015-10-17 15:45:17,305 INFO [main] org.apache.hadoop.mapred.MapTask: Spilling map output
2015-10-17 15:45:17,305 INFO [main] org.apache.hadoop.mapred.MapTask: bufstart = 0; bufend = 48271024; bufvoid = 104857600
2015-10-17 15:45:17,305 INFO [main] org.apache.hadoop.mapred.MapTask: kvstart = 26214396(104857584); kvend = 17310640(69242560); length = 890375$
2015-10-17 15:45:17,305 INFO [main] org.apache.hadoop.mapred.MapTask: (EQUATOR) 57339776 kvi 14334940(57339760)
2015-10-17 15:45:26,696 INFO [SpillThread] org.apache.hadoop.mapred.MapTask: Finished spill 0
2015-10-17 15:45:26,696 INFO [main] org.apache.hadoop.mapred.MapTask: (RESET) equator 57339776 kv 14334940(57339760) kvi 12140764(48563056)
2015-10-17 15:45:30,603 INFO [main] org.apache.hadoop.mapred.MapTask: Spilling map output
2015-10-17 15:45:30,603 INFO [main] org.apache.hadoop.mapred.MapTask: bufstart = 57339776; bufend = 743078; bufvoid = 104857600
2015-10-17 15:45:30,603 INFO [main] org.apache.hadoop.mapred.MapTask: kvstart = 14334940(57339760); kvend = 5428644(21714576); length = 8906297/$
2015-10-17 15:45:30,603 INFO [main] org.apache.hadoop.mapred.MapTask: (EQUATOR) 9811814 kvi 2452948(9811792)
2015-10-17 15:45:39,525 INFO [SpillThread] org.apache.hadoop.mapred.MapTask: Finished spill 1
2015-10-17 15:45:39,525 INFO [main] org.apache.hadoop.mapred.MapTask: (RESET) equator 9811814 kv 2452948(9811792) kvi 244148(976592)
2015-10-17 15:45:43,307 INFO [main] org.apache.hadoop.mapred.MapTask: Spilling map output
2015-10-17 15:45:43,307 INFO [main] org.apache.hadoop.mapred.MapTask: bufstart = 9811814; bufend = 58036090; bufvoid = 104857600
2015-10-17 15:45:43,307 INFO [main] org.apache.hadoop.mapred.MapTask: kvstart = 2452948(9811792); kvend = 19751904(79007616); length = 8915445/6$
2015-10-17 15:45:43,307 INFO [main] org.apache.hadoop.mapred.MapTask: (EQUATOR) 67104842 kvi 16776204(67104816)
```


Dump

Dumping is the same as logging, but it is usually a once-in-a-while occurrence because the data to write is huge

Things to dump

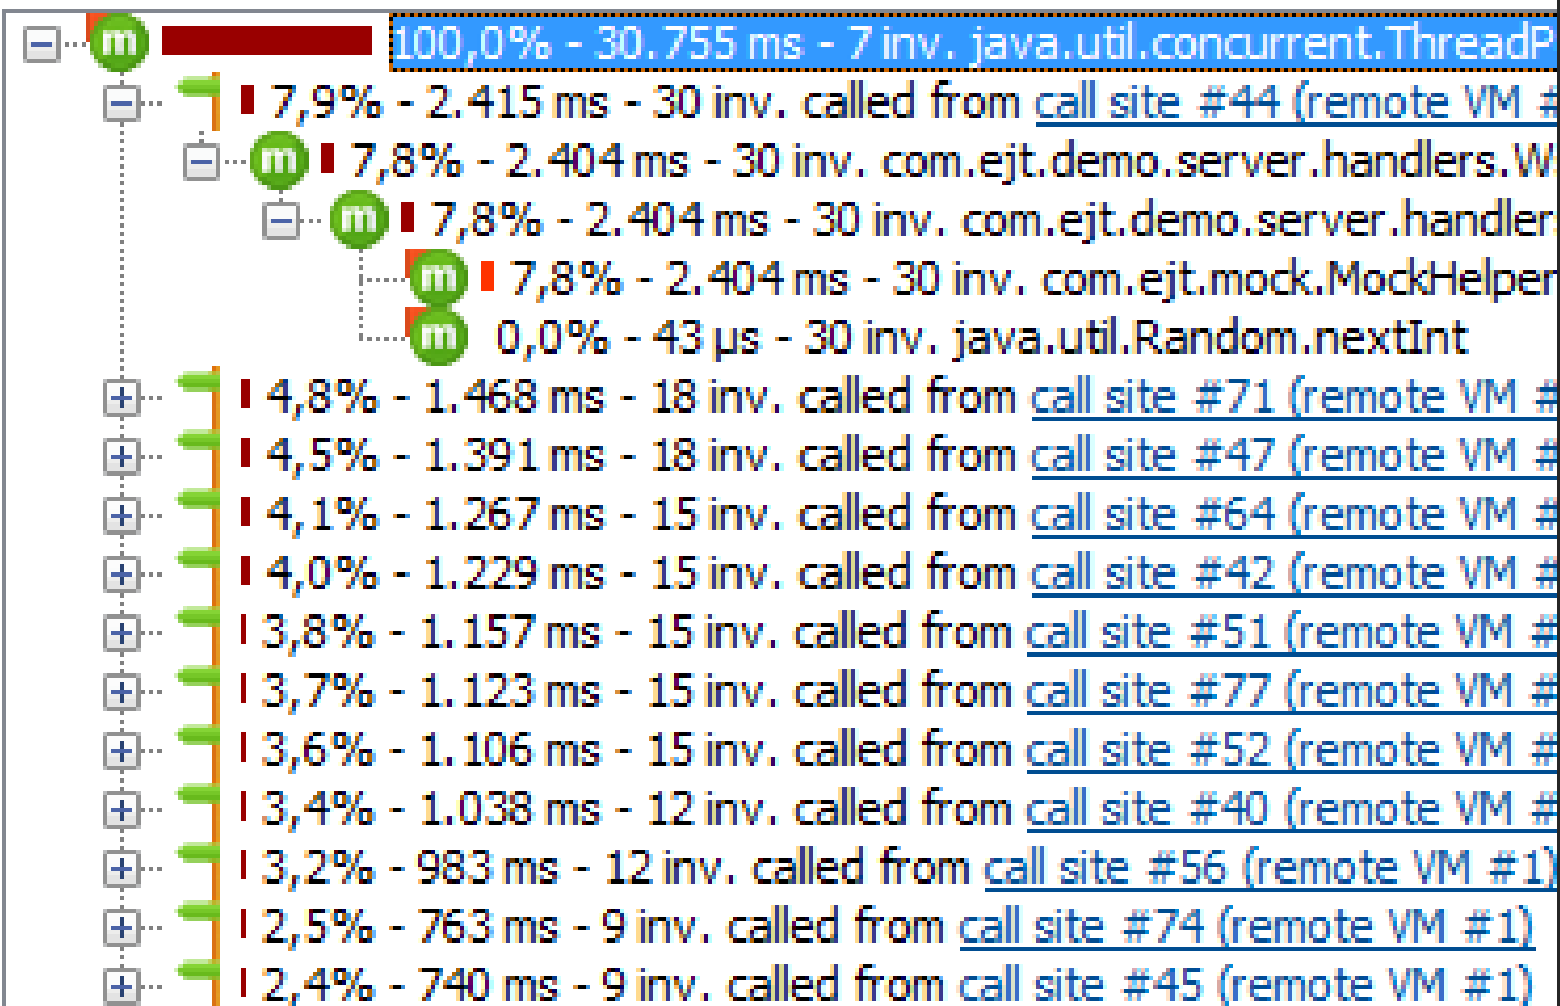
- Memory
- Databases
- Stack Traces

C:\Windows\System32\cmd.exe - testrom -b test.bin

> dump \$2000 \$2200

```
03-2000: 4083 130D 1311 1314 021F 2023 4C49 5354 0..... #LIST
03-2008: 204B 4559 A807 6C0F 0000 0828 47D1 443A KEY..1....<G.D:
03-2010: 44E0 0360 F187 607C 2034 321D 07EC 0A1D D...`...! 42.....
03-2018: 47D7 F082 44F6 6C05 6D2D 4526 E8C6 F081 G...D.l.m-E&....
03-2020: 4523 F081 F081 2060 4C49 5354 A805 6C2C E#.... `LIST..1,
03-2028: 0000 0A34 C12C 082A 47B3 441C 7171 834B ...4...*G.D.qq.K
03-2030: 2025 B800 7171 834B 2026 321D 0001 B21D %..qq.K &2.....
03-2038: 834B 2025 8000 F020 2001 F403 47B8 6FBC .K %... ..G.o.
03-2040: 0360 F187 607C 2038 321D 07BB 0A1D 47A8 ...`...! 82.....G.
03-2048: F082 7170 3360 7171 834B 2023 B800 7171 ..qp3`qq.K #..qq
03-2050: 2021 B800 F081 44BD 6C09 6D03 C103 E8C6 ?....D.l.m.....
03-2058: 4504 6D01 C103 E8C6 F081 C103 6CFA 6FF9 E.m.....l.o.
03-2060: 2077 5255 4E20 A805 6C0F 0000 4406 F081 wRUN ..1...D...
03-2068: 0796 086C 4785 F082 0824 4772 7171 834B ...lG....$Grqq.K
03-2070: 2023 B800 F082 C0CE C103 6CE3 F081 208A #.....l...
03-2078: 434F 4E54 A805 6C08 0000 47EF F081 077E CONT..1...G....~
03-2080: 086C 476E F082 C0CE C103 6C02 F081 0887 .lGn.....l.....
03-2088: C0D3 F081 2097 5343 5241 5443 4820 4120 .... .SCRATCH A
03-2090: A808 E8CE 0000 C168 0000 0087 6C3E 20AA .....h....l> .
03-2098: 5343 5241 5443 4820 4B45 5920 A809 6C07 SCRATCH KEY ..1.
03-20A0: 0000 444D 0761 0834 474B F082 C0CE 449C ..DM.a.4GK....D.
03-20A8: F081 F081 20B5 5343 5241 5443 4820 5020 .... .SCRATCH P
03-20B0: A808 E8CE 0000 0022 6C20 20C0 5343 5241 ..... "1 .SCRA
```

Press STOP to quit, B to step back or any other key to proceed_



As you can imagine, **having archives** is a **BIG** part of **operations and maintenance** teams to find and solve bugs **caused by users** that your team did not catch.

App
crashed on
user's side



We
have logs



We didn't
log the part
of the code
that crashed



Remember

Keeping an archive is **not free**.

Files **accumulate** and you need to manage them somehow.

Problems that you might eventually deal with

Logging to gather data

Log rotation for apps that run 24/7

Log archiving

Log merging (parallel programs)

Log filtering for different teams

Crash logs from users



‘Debuggers’

Debuggers are programs that attaches itself to a running program and provide various information to you as the program runs.

Notable debuggers for C/C++

Visual Studios

GDB

Valgrind

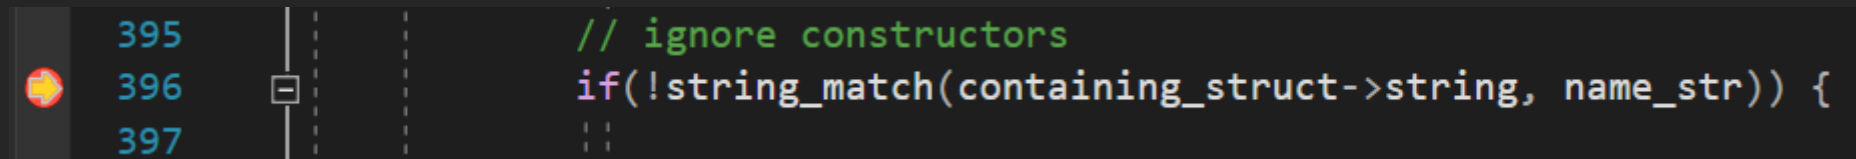
WinDBG

RemebyBG

and many many more...

Notable debugger features

Most debuggers come with **breakpoints**.
They allow you to ‘pause’ your program while
it’s running when a line is ‘hit’.



A screenshot of a debugger's source code window. On the left, line numbers 395, 396, and 397 are listed. A red pushpin icon is next to line 396, indicating a breakpoint. A small square icon is also present next to line 396. The code for line 396 is `if(!string_match(containing_struct->string, name_str)) {`. The code for line 395 is `// ignore constructors`. The code for line 397 is partially visible as `...`.

```
395 // ignore constructors
396 if(!string_match(containing_struct->string, name_str)) {
397
```

When a **breakpoint** is hit, there are a bunch of things you can inspect/edit including variables, memory, disassembly

Watches


Watch 1

Search (Ctrl+E) Search Depth: 3

Name	Value	Type
n	1	int
f	-107374176.	float
Add item to watch		

Autos Locals Watch 1

Call Stack

Call Stack	
	Name
	Project1.exe!fib(int n) Line 6
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!fib(int n) Line 9
	Project1.exe!main() Line 13
	[External Code]
	kernel32.dll

Memory

Memory 2

Address: 0x000000EED8FAFBB2

0x000000EED8FAFBB2	c3	ce	f7	7f	00	00	31	00	ÃÎ÷...1.
0x000000EED8FAFBBA	5c	00	78	00	36	00	34	00	\.x.6.4.
0x000000EED8FAFBC2	5c	00	44	00	65	00	62	00	\.D.e.b.
0x000000EED8FAFBCA	75	00	67	00	5c	00	e6	cd	u.g.\.æÍ
0x000000EED8FAFBD2	8e	30	e3	8a	59	f7	65	00	Ž0ăŠY÷e.
0x000000EED8FAFBDA	63	00	74	00	31	00	2e	00	c.t.1...
0x000000EED8FAFBE2	65	00	78	00	65	00	96	55	e.x.e.-U
0x000000EED8FAFBEA	90	17	fb	7f	00	00	ff	00	..û...ÿ.
0x000000EED8FAFBF2	00	40	31	65	59	f7	19	00	..@1eY÷..
0x000000EED8FAFBFA	00	00	fb	7f	00	00	ff	00	..û...ÿ.
0x000000EED8FAFC02	00	40	31	65	59	f7	19	00	..@1eY÷..
0x000000EED8FAFC0A	00	00	fb	7f	00	00	68	48	..û...hH
0x000000EED8FAFC12	a1	17	fb	7f	00	00	50	87	j.û...P.
0x000000EED8FAFC1A	90	17	fb	7f	00	00	e6	cd	..û...æÍ
0x000000EED8FAFC22	8e	30	e3	8a	59	f7	50	87	Ž0ăŠY÷P.
0x000000EED8FAFC2A	90	17	fb	7f	00	00	00	10	..û.....
0x000000EED8FAFC32	c3	ce	f7	7f	00	00	00	00	ÃÎ÷.....
0x000000EED8FAFC3A	00	00	00	00	00	00	00	00
0x000000EED8FAFC42	00	00	00	00	00	00	3b	18;.
0x000000EED8FAFC4A	c2	ce	f7	7f	00	00	00	10	ÃÎ÷.....
0x000000EED8FAFC52	c3	ce	f7	7f	00	00	a0	98	ÃÎ÷... ~
0x000000EED8FAFC5A	b2	ac	fb	7f	00	00	4c	9c	..¬û...Lœ
0x000000EED8FAFC62	95	74	ca	02	00	00	a0	98	..tÊ... ~
0x000000EED8FAFC6A	b2	ac	fb	7f	00	00	02	00	..¬û.....
0x000000EED8FAFC72	00	00	ee	00	00	00	bb	98	..î...»~
0x000000EED8FAFC7A	b2	ac	fb	7f	00	00	05	00	..¬û.....

Disassembly

```
91     current_easing_function_index = 0;
92 }
93 void GameUpdate(void)
94 {
95     timer += CP_System_GetDt();
96     if (timer >= duration) {
97         timer = 0.f;
98         SWAP(float, min_x, max_x);
99         SWAP(float, min_size, max_size);
100     }
101
102
103     if (CP_Input_KeyTriggered(KEY_RIGHT))
104     {
105         ++current_easing_function_index;
106         if (current_easing_function_index >= ArrayCount(ea
```

Viewing Options

```
00007FF7C4DE3E17 call    __CheckForDebuggerJustMyCode (07FF7C4DE139Dh)
        timer += CP_System_GetDt();
00007FF7C4DE3E1C call    qword ptr [__imp_CP_System_GetDt (07FF7C4DF4000h)]
00007FF7C4DE3E22 movss   xmm1,dword ptr [timer (07FF7C4DF0B44h)]
00007FF7C4DE3E2A addss   xmm1,xmm0
00007FF7C4DE3E2E movaps  xmm0,xmm1
00007FF7C4DE3E31 movss   dword ptr [timer (07FF7C4DF0B44h)],xmm0
        if (timer >= duration) {
00007FF7C4DE3E39 movss   xmm0,dword ptr [timer (07FF7C4DF0B44h)]
00007FF7C4DE3E41 comiss  xmm0,dword ptr [duration (07FF7C4DF0B40h)]
00007FF7C4DE3E48 jnb     __$EncStackInitStart+0AAh (07FF7C4DE3EA9h)
        timer = 0.f;
00007FF7C4DE3E4A xorps   xmm0,xmm0
00007FF7C4DE3E4D movss   dword ptr [timer (07FF7C4DF0B44h)],xmm0
        SWAP(float, min_x, max_x);
```

or you can also use <https://godbolt.org/>

Visual Studios Debugger Demo

Internal debugger

Instead of using an external debugger at hooks onto your program, you can write your own a debugger **WITHIN** your program.

The good thing is that the variables and functions you want to debug are **immediately available to you** and that it will be **targeted to best fit your program.**

Minecraft 1.17.1 (fabric-loader-0.11.6-1.17.1/fabric/Fabric)

Java: 16.0.1 64bit

56 fps T: inf vsync fancy fancy-clouds B: 2

Mem: 38% 782/2048MB

"vanilla" server, 2 tx, 1432 rx

Allocated: 78% 1600MB

C: 231/4624 (s) D: 8, pC: 000, pU: 00, aB: 08

CPU: 8x Intel(R) Core(TM) i7-6700 CPU @ 3.40GHz

E: 9/389 B: 0

F: 220 T: 389

Chunks[C] W: 625, 420 E: 389,139,420

Display: 1920x1057 (NVIDIA Corporation)

minecraft:overworld FC: 0

GeForce GTX 1070/PCIe/SSE2

295, 65, -17

3.2.0 NVIDIA 456.71

XYZ: 295.512 / 65.00000 / -16.827

Block: 295 65 -17

Chunk: 7 1 15 in 18 4 -2

Facing: east (Towards positive X) (-99.3 / 4.8)

Client Light: 15 (15 sky, 10 block)

CH S: 64 M: 64

SH S: 64 O: -1 M: 64 ML: -1

Biome: minecraft:wooded_hills

Local Difficulty: 2.25 / / 0.13 (Day 1774)

Sounds: 17/247 + 1/8 (Mood 0%)

[Fabric] Active renderer: IndigoRenderer

Debug: Pie [shift]: hidden FPS [alt]: hidden

For help: press F3 + O



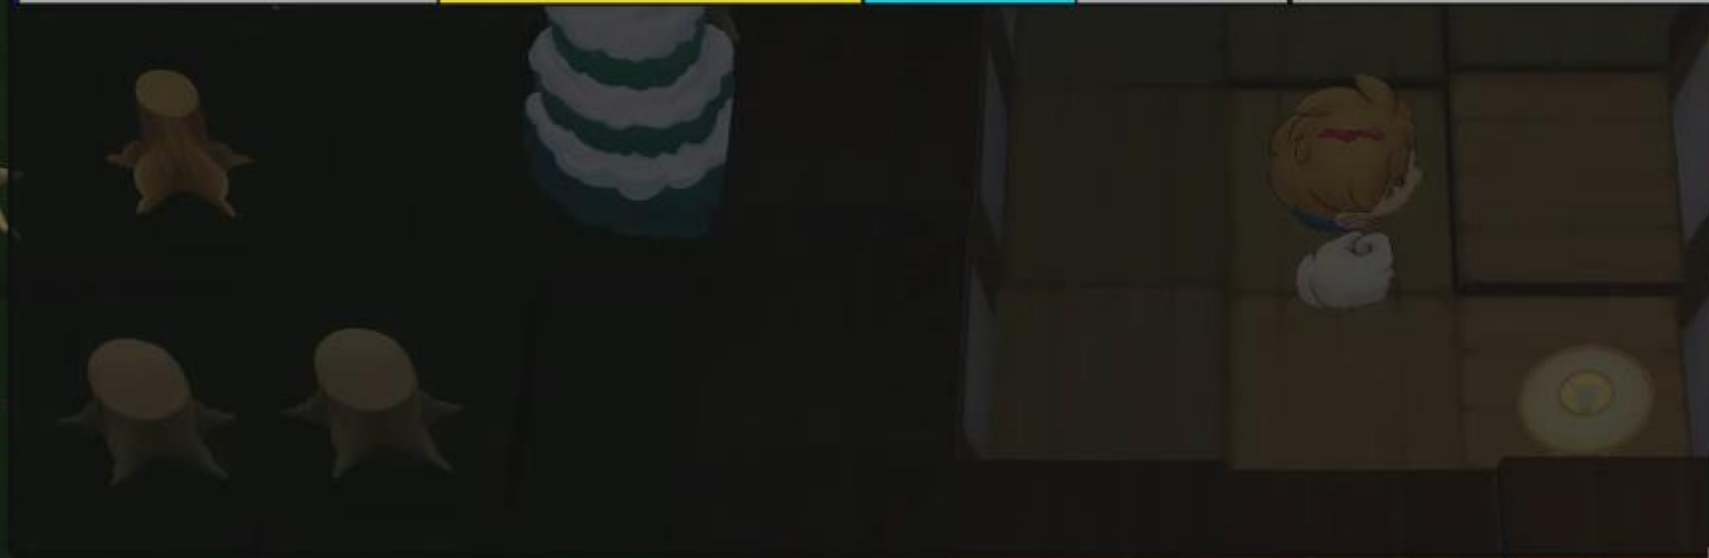
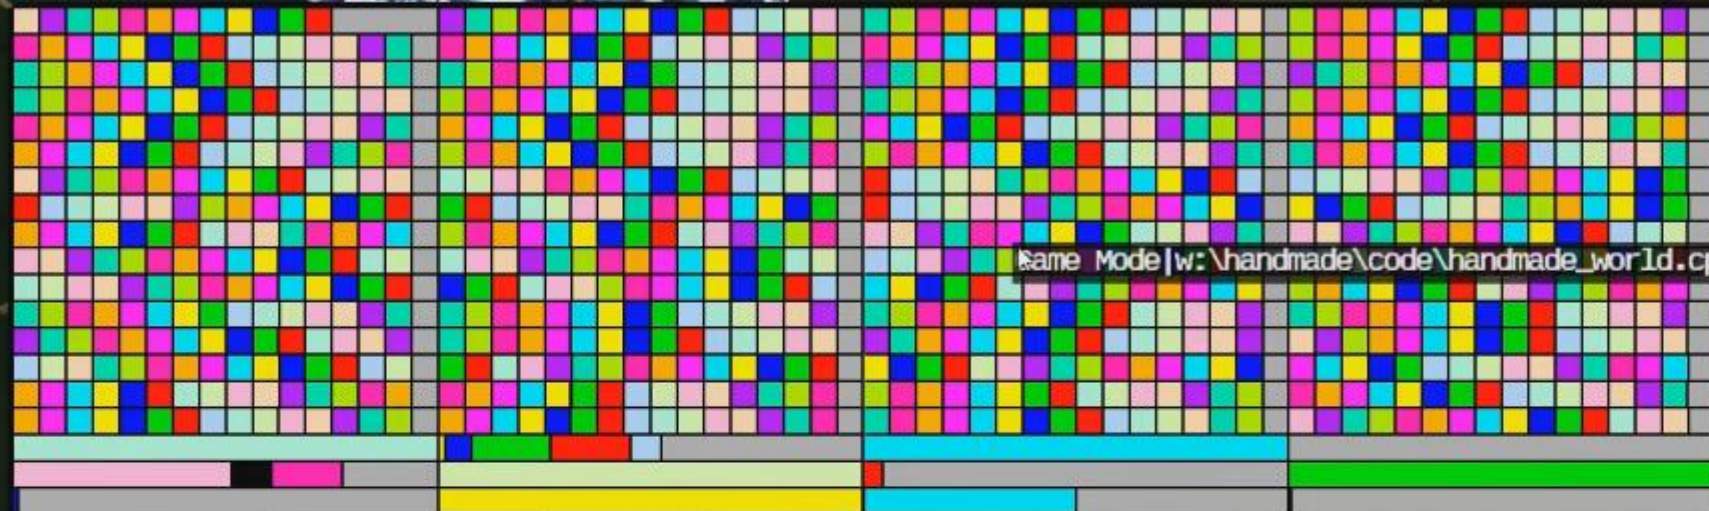
8



Zombie dies
Footsteps
Villager mumbles

Memory

Arenas Frames Sizes Debug Call Sites



help commands

bind: Creates, deletes or lists key bindings associated to commands or shortcuts.
clear: Clears the in-game console.
help: Gets help on using commands.

Listed 3/3 commands.

bind

Listed 0/0 key bindings.

> _

At this point, how you want to debug
is up to your **creativity!**

At the end of the day, debugging is **part of the engineering process**. For every route you can take, you must **think critically**!

Also, don't forget talking

Sometimes, it might be helpful to speak what
you are thinking to another person.
(sounds ridiculous, but it might work for some)

