

# OS Architecture and System Calls

Instructor: William Zheng

Email:

[william.zheng@digipen.edu](mailto:william.zheng@digipen.edu)

PHONE EXT: 1745

# A word about the study of OS

- Generalist versus Specialist
- Organic and circularity

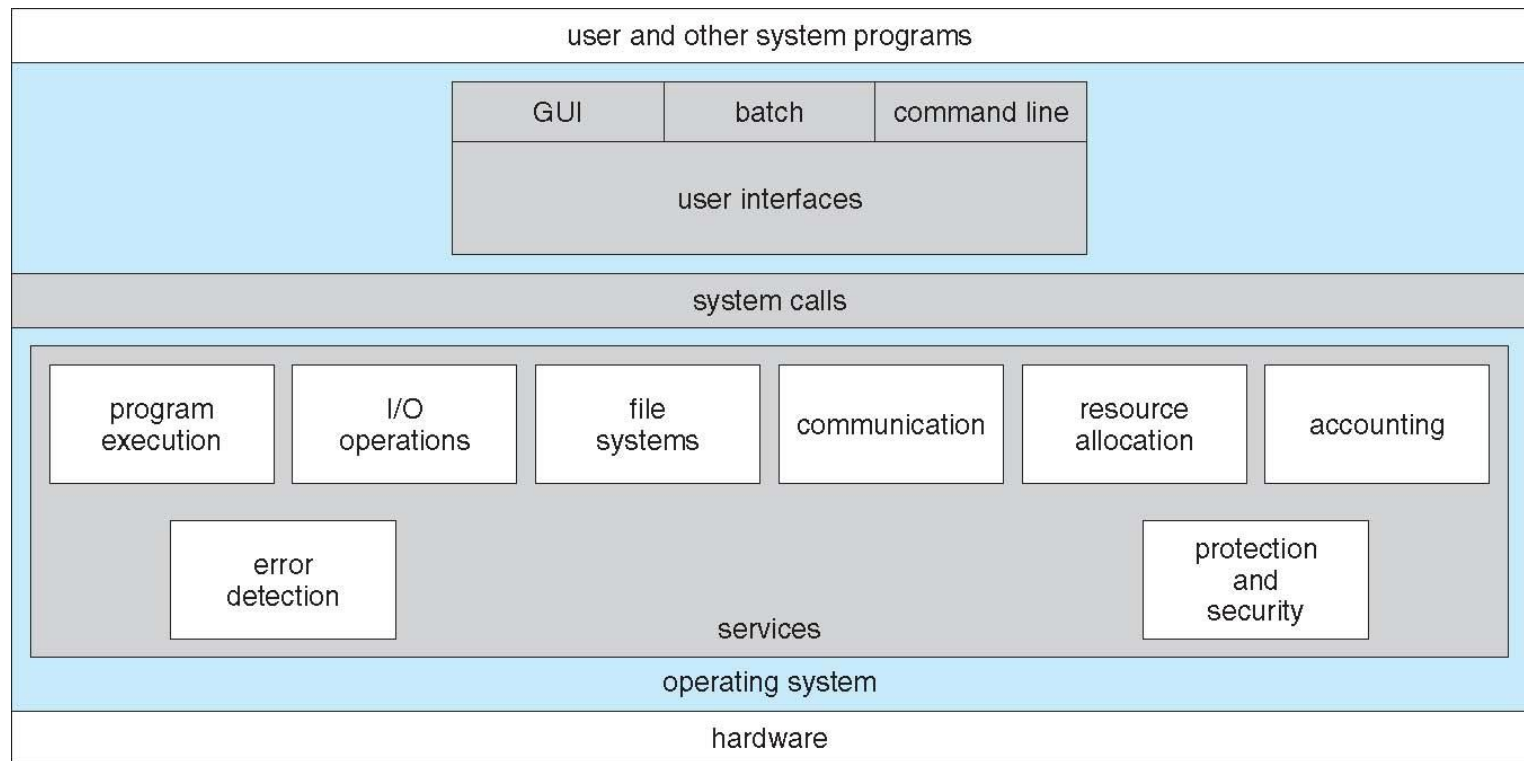
# Goals of this lecture

- System Calls
  - Interface between user program and OS
- OS architecture
  - Organization and Design of OS code

# What is an OS and what are the roles of OS?

- Process Management
  - A process a running program.
  - Starting, terminating processes
  - Coordinating processes (allowing processes to talk to one another, protecting them from one another)
- Memory Management
  - Allocation of memory for the processes
  - Protect memory from being written by another process
  - How to handle when we run out of RAM
- I/O Management
  - Device drivers
  - Providing API for the user programs (encapsulation principle)
- Storage Management
  - File systems and data

# OS Services



# Invoking the OS services

- **int** instruction
  - Generate the software interrupt
    - **int 80h**
      - 80h: the interrupt vector for system call
  - Return: `iret`
- **sysenter/sysexit** instruction
- **syscall** instruction

Main:

```
...  
call FUNC  
...  
ret
```

FUNC:

```
...  
...call the OS. How?  
...  
ret
```

# System calls

- API provided by OS
  - For access to services provided by OS
- How to access?
  - software-generated interrupt (usually)
  - **call** instruction (older days)

# Why system calls necessary

## ***Access to privileged operations:***

- *Many operations, such as managing hardware devices or modifying system configurations, require higher privileges that are only accessible through system calls.*

## ***Resource management:***

- *System calls provide a standardized interface for allocating and managing system resources like memory, files, and devices, ensuring fair and controlled access by different processes.*



# Why system calls necessary

## ***Abstraction:***

- *System calls abstract the underlying complexities of the operating system, allowing application developers to interact with the system in a higher-level, platform-independent manner.*

## ***Security and protection:***

- *System calls enforce access control and security policies, preventing unauthorized access to sensitive resources and protecting the integrity of the system.*

# Types of system calls – Windows/Linux

Process	Windows	Linux
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	Fork() Exit() Wait()
File manipulation	CreateFile() ReadFile() WriteFile()	Open() Read() Write() Close()
Device Management	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() Read() Write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	Getpid() Alarm() Sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	Pipe() Shmget() Mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	Chmod() Umask() Chown()

# System call example

- Copy contents of file A to file B
  - How many system calls involved?

# Difference between system call and function call

- System call: a call into kernel code, typically performed by executing an interrupt
- Function call, if calling a system call (via library), switches into kernel mode from user mode

# Portability - I

Program X



Computer  
System A

Program X

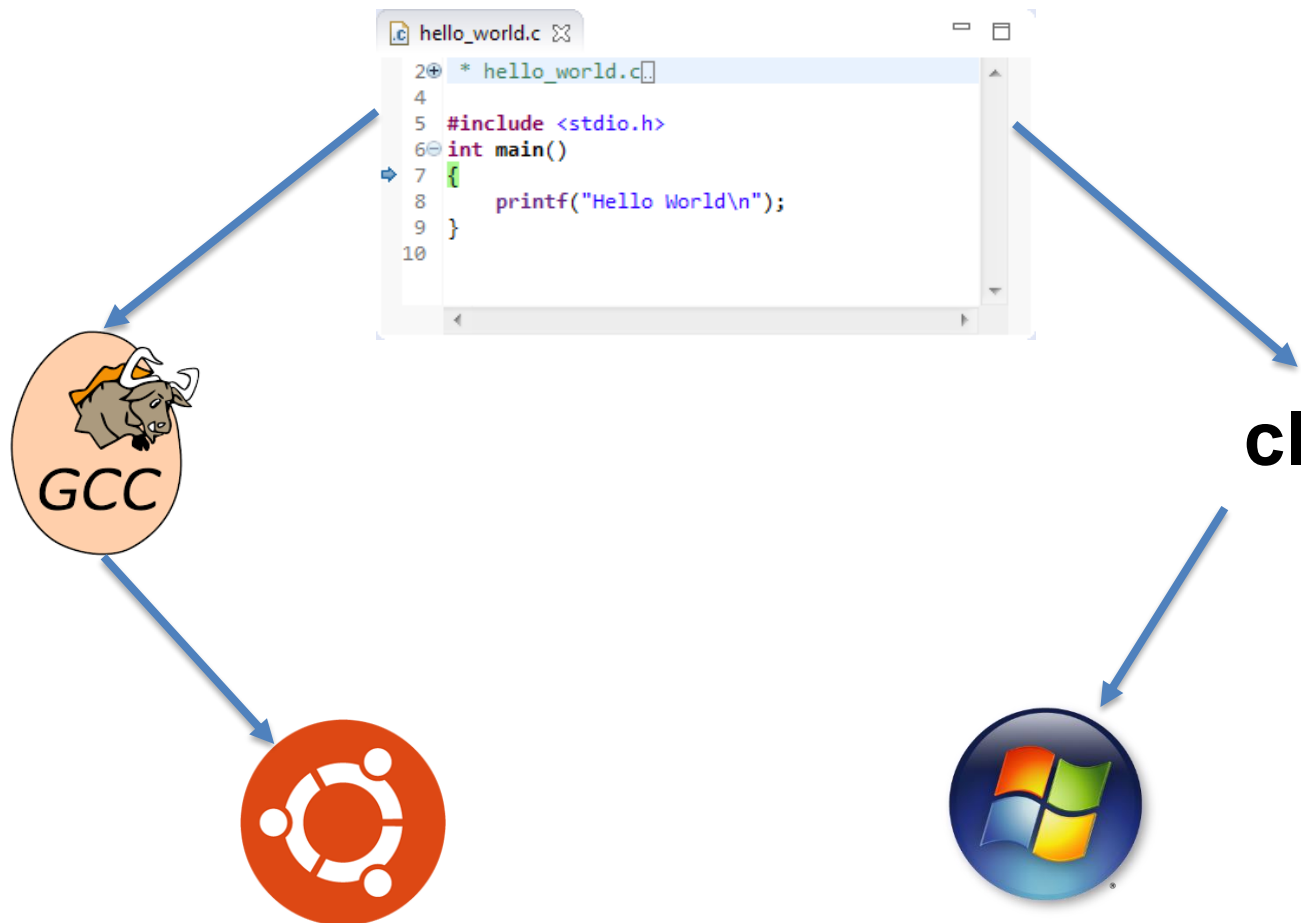


Computer  
System B

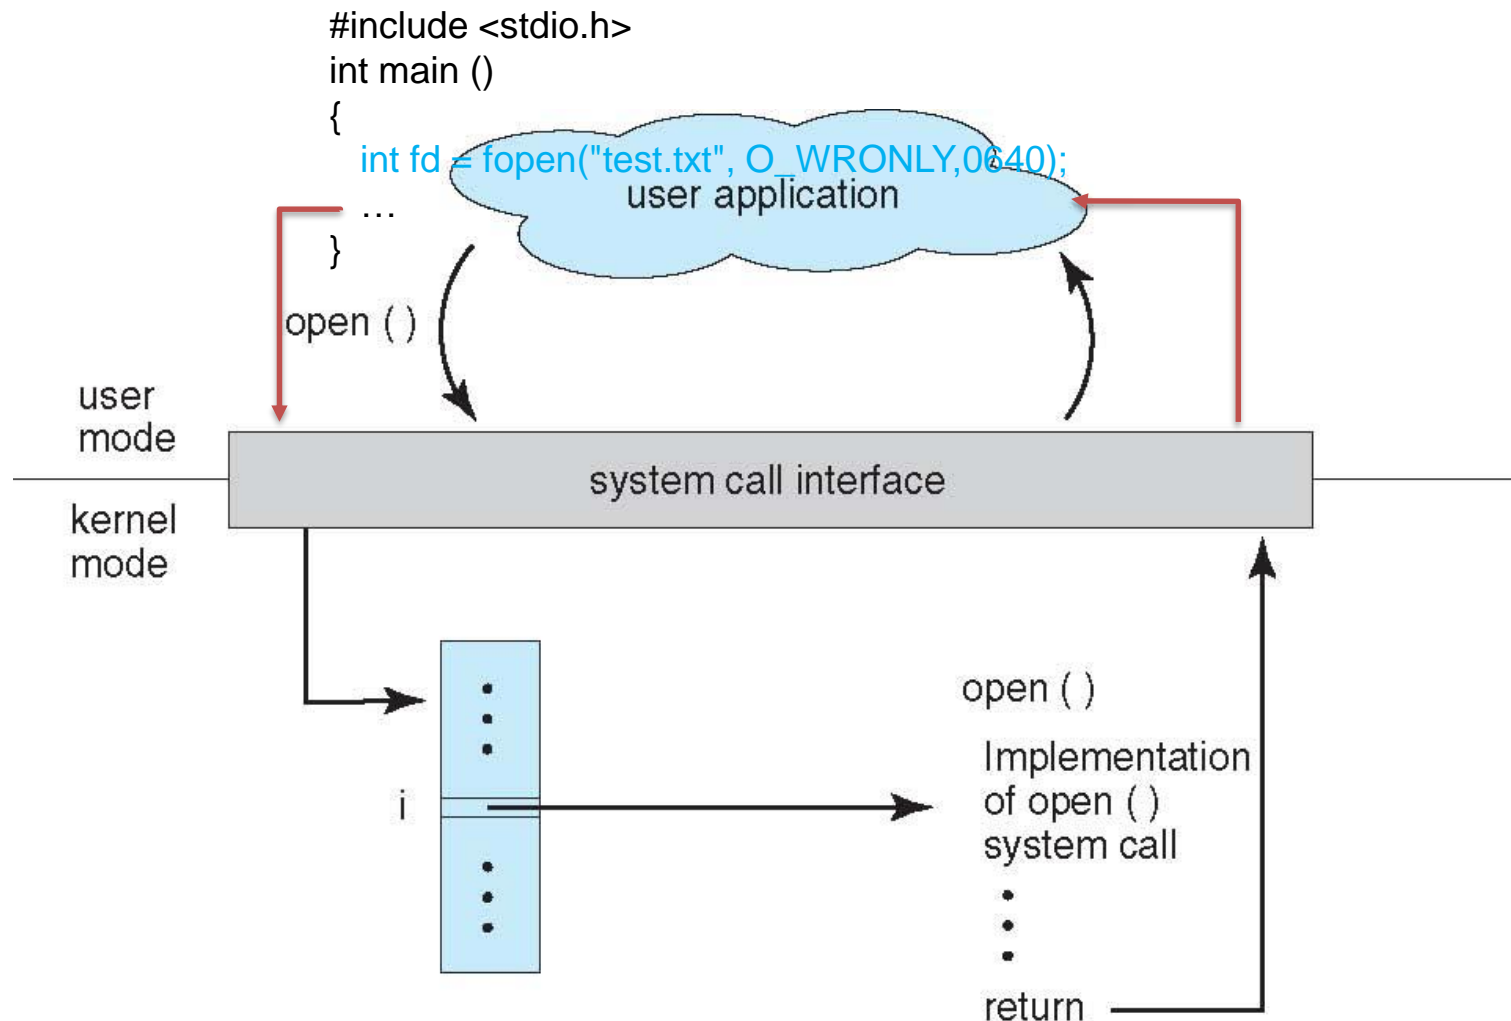
# Portability - II

- Binary level
- Source code level

# Source Code Portability - Demo



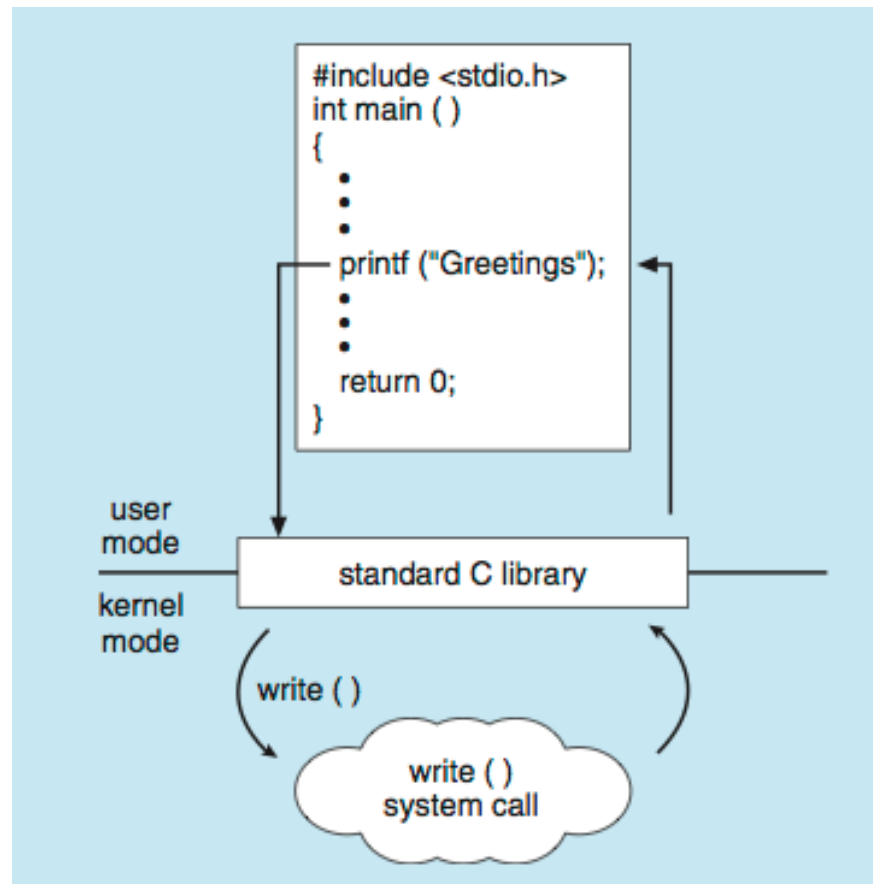
# System calls, API and portability



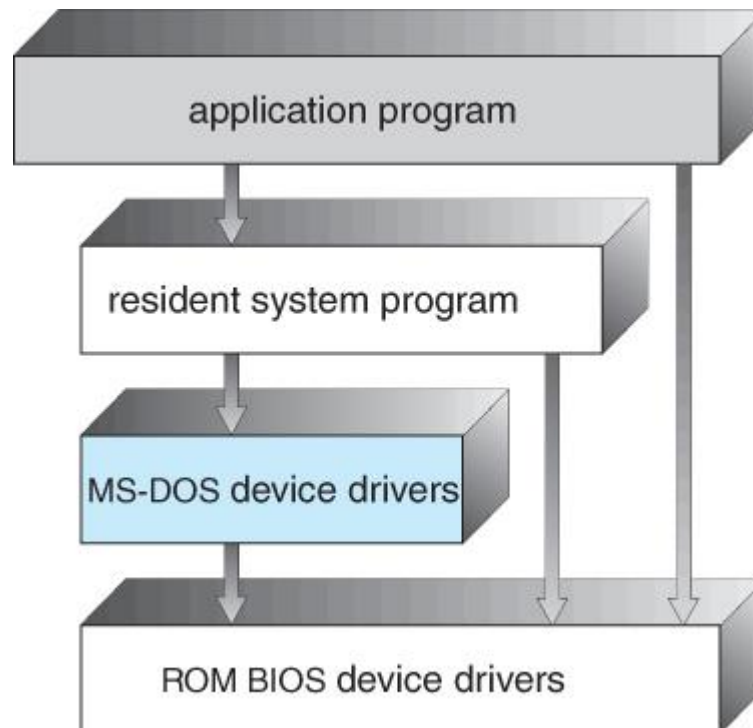


# Standard C Library Example

C program invoking printf() library call, which calls write() system call




# DOS OS architecture - I




# DOS OS architecture - II

- Single-tasking
- Shell invoked when system booted
- Simple method to run program
  - No process created
- Single memory space
- Loads program into memory, overwriting all but the kernel
- Program exit -> shell reloaded




(a)



(b)

# OS architecture

- Everything \*almost\* runs in kernel model
- Example: Linux
- Pros and Cons



# Monolithic OS architecture

**Pros:** Good performance (**S**peed) / (Shared kernel space) Simplicity of design

**Cons:**

No information hiding

Potential stability issues / Chaotic

Hard to understand

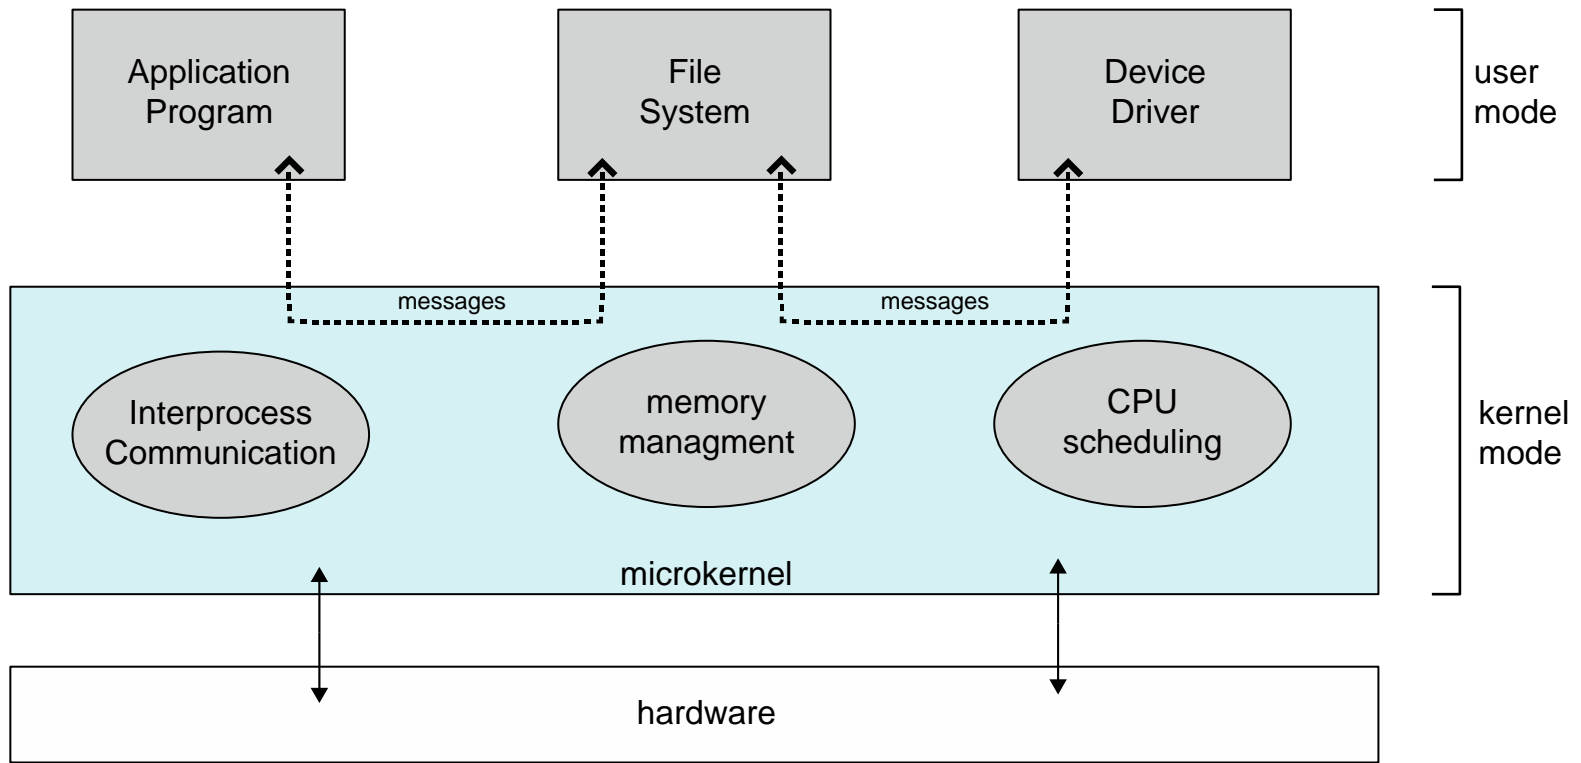
- Can become huge
- Linux 4.15 has 20 million lines of code
- Windows 10 contains over 40 million lines!
- Potentially difficult to maintain

**Examples:**

- Traditional Unix kernels (includes BSDs and Solaris)
- Linux
- MS-DOS, Windows 9x
- Mac OS versions below 8.6

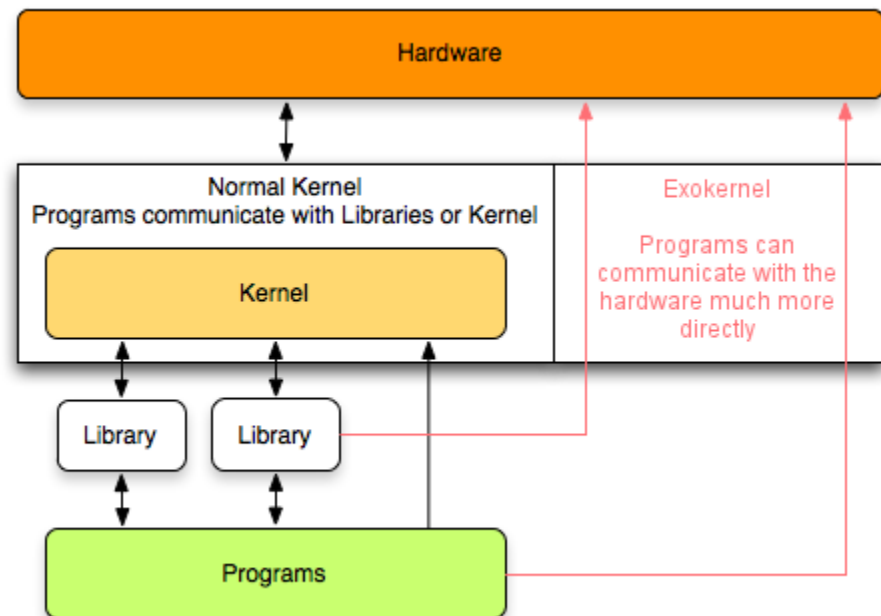
# Microkernel architecture

- Push stuffs out of kernel
- Client-Server model
- Example: Mac OS, iPhone OS etc
- Pros : Modularity: easier management, Fault isolation and reliability
- Cons: Inefficient (boundary crossings); Insufficient protection; Inconvenient to share data between kernel and services

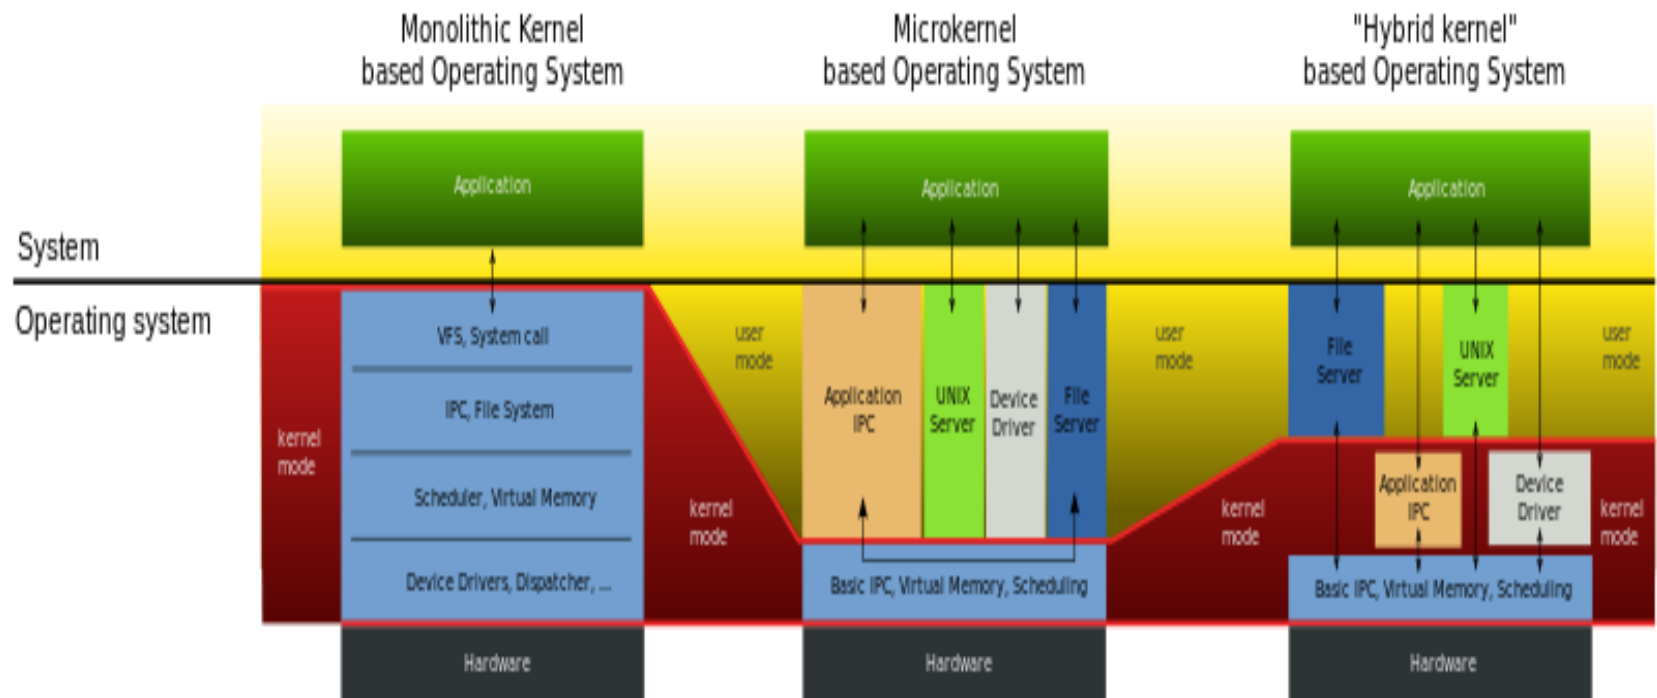


# Exokernel architecture

- Proposed by MIT
- Application-oriented OS
  - Gives them all control
- Research



# Comparison





# Questions

- Is the operating system using the CPU at all times?
- After booting and initialization, what are the circumstances that would cause OS code to use the CPU?