

# Operating System: ELF Format

July 29, 2023

## 1 ELF format header

1. Compile `hello-world.cpp` under WSL.
2. Use the command `file a.out`, where `a.out` is the compiled output executable.

```
a.out: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV), dynamically linked,
interpreter /lib64/ld-linux-x86-64.so.2,
BuildID[sha1]=2317c609af99bb13a86e3c299af29c9a650777e3,
for GNU/Linux 3.2.0, not stripped
```

3. Use the command `readelf -h a.out` to read the header of ELF format, which is defined as follows<sup>1</sup>

```
# define EI_NIDENT 16

typedef struct {
    unsigned char
    e_ident[EI_NIDENT];
    Elf64_Half e_type;
    Elf64_Half e_machine;
    Elf64_Word e_version;
    Elf64_Addr e_entry;
    Elf64_Off e_phoff;
    Elf64_Off e_shoff;
    Elf64_Word e_flags;
    Elf64_Half e_ehsize;
    Elf64_Half e_phentsize;
    Elf64_Half e_phnum;
    Elf64_Half e_shentsize;
    Elf64_Half e_shnum;
    Elf64_Half e_shstrndx;
} Elf64_Ehdr;
```

The ELF header information of `a.out` is :

---

<sup>1</sup>Please refer to here.

#### ELF Header:

```
Magic:  7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
Class:                                     ELF64
Data:                                     2's complement, little endian
Version:                                 1 (current)
OS/ABI:                                 UNIX - System V
ABI Version:                             0
Type:                                    DYN (Position-Independent Executable file)
Machine:                                Advanced Micro Devices X86-64
Version:                                0x1
Entry point address:                     0x1100
Start of program headers:                 64 (bytes into file)
Start of section headers:                 14192 (bytes into file)
Flags:                                    0x0
Size of this header:                      64 (bytes)
Size of program headers:                   56 (bytes)
Number of program headers:                 13
Size of section headers:                   64 (bytes)
Number of section headers:                 31
Section header string table index: 30
```

## 2 ELF section

1. Use the following command `readelf -S a.out` to see all the sections that are present in the example `hello-world.cpp` :

```
1 There are 31 section headers, starting at offset 0x3770:
3 Section Headers:
   [Nr] Name                Type                Address                Offset
   [ 0] Size                EntSize              Flags  Link  Info  Align
7   [ 0] 0000000000000000 NULL              0000000000000000      0  0  0
9   [ 1] .interp                 PROGBITS          0000000000000318    00000318
   [ 1] 000000000000001c        0000000000000000      A  0  0  1
11  [ 2] .note.gnu.pr[...]      NOTE              0000000000000338    00000338
   [ 2] 0000000000000030        0000000000000000      A  0  0  8
13  [ 3] .note.gnu.bu[...]      NOTE              0000000000000368    00000368
   [ 3] 0000000000000024        0000000000000000      A  0  0  4
15  [ 4] .note.ABI-tag          NOTE              000000000000038c    0000038c
   [ 4] 0000000000000020        0000000000000000      A  0  0  4
17  [ 5] .gnu.hash              GNU_HASH          00000000000003b0    000003b0
   [ 5] 0000000000000024        0000000000000000      A  6  0  8
19  [ 6] .dynsym                DYNAMIC           00000000000003d8    000003d8
   [ 6] 0000000000000120        0000000000000018      A  7  1  8
21  [ 7] .dynstr                STRTAB            00000000000004f8    000004f8
   [ 7] 00000000000000be        0000000000000000      A  0  0  1
23  [ 8] .gnu.version            VERSYM            00000000000005b6    000005b6
   [ 8] 0000000000000018        0000000000000002      A  6  0  2
25  [ 9] .gnu.version_r          VERNEED           00000000000005d0    000005d0
   [ 9] 0000000000000040        0000000000000000      A  7  1  8
27  [10] .rela.dyn              RELA              0000000000000610    00000610
   [10] 00000000000000c0        0000000000000018      A  6  0  8
```

29	[11]	.rela.plt	RELA	000000000000006d0	000006d0
		00000000000000090	0000000000000018	AI 6 24	8
	[12]	.init	PROGBITS	00000000000001000	00001000
31		0000000000000001b	0000000000000000	AX 0 0	4
	[13]	.plt	PROGBITS	00000000000001020	00001020
33		00000000000000070	0000000000000010	AX 0 0	16
	[14]	.plt.got	PROGBITS	00000000000001090	00001090
35		00000000000000010	0000000000000010	AX 0 0	16
	[15]	.plt.sec	PROGBITS	000000000000010a0	000010a0
37		00000000000000060	0000000000000010	AX 0 0	16
	[16]	.text	PROGBITS	00000000000001100	00001100
39		000000000000001bd	0000000000000000	AX 0 0	16
	[17]	.fini	PROGBITS	000000000000012c0	000012c0
41		0000000000000000d	0000000000000000	AX 0 0	4
	[18]	.rodata	PROGBITS	00000000000002000	00002000
43		0000000000000000a	0000000000000000	A 0 0	4
	[19]	.eh_frame_hdr	PROGBITS	0000000000000200c	0000200c
45		00000000000000034	0000000000000000	A 0 0	4
	[20]	.eh_frame	PROGBITS	00000000000002040	00002040
47		0000000000000000ac	0000000000000000	A 0 0	8
	[21]	.init_array	INIT_ARRAY	00000000000003d90	00002d90
49		000000000000000008	0000000000000008	WA 0 0	8
	[22]	.fini_array	FINLARRAY	00000000000003d98	00002d98
51		000000000000000008	0000000000000008	WA 0 0	8
	[23]	.dynamic	DYNAMIC	00000000000003da0	00002da0
53		0000000000000001f0	0000000000000010	WA 7 0	8
	[24]	.got	PROGBITS	00000000000003f90	00002f90
55		00000000000000070	0000000000000008	WA 0 0	8
	[25]	.data	PROGBITS	00000000000004000	00003000
57		00000000000000010	0000000000000000	WA 0 0	8
	[26]	.bss	NOBITS	00000000000004010	00003010
59		000000000000000008	0000000000000000	WA 0 0	1
	[27]	.comment	PROGBITS	00000000000000000	00003010
61		0000000000000002b	0000000000000001	MS 0 0	1
	[28]	.symtab	SYMTAB	00000000000000000	00003040
63		000000000000003d8	0000000000000018	29 18	8
	[29]	.strtab	STRTAB	00000000000000000	00003418
65		0000000000000023d	0000000000000000	0 0	1
	[30]	.shstrtab	STRTAB	00000000000000000	00003655
67		0000000000000011a	0000000000000000	0 0	1
Key to Flags:					
69	W (write), A (alloc), X (execute), M (merge), S (strings), I (info),				
	L (link order), O (extra OS processing required), G (group), T (TLS),				
71	C (compressed), x (unknown), o (OS specific), E (exclude),				
	D (mbind), l (large), p (processor specific)				

sections.txt

From the output, it can be seen that there are **.text**, **.data**, **.bss**, **.symtab** and **.shstrtab** sections.

2. To check the code that is present in **.text**, use the command `objdump -d a.out`.
3. How does the compiler come to know where to put each section in the final ELF. Linkers use a file called the linker descriptor file. It contains information about all the memory in the target machine with its starting address and size, as well as information about the different sections that should be present in the final ELF file and where each section should be loaded in the target machine.

4. The program header in the ELF specifies which section of the ELF file goes to a particular memory location. The program header location in the ELF file is provided by the `e_phoff` variable present in the ELF header. `e_phoff` provides the offset at which the ELF program header is present in the ELF file. `e_phnum` is the number of program header entries in the ELF file and the size of each entry is `e_phentsize`. Use the following command to see the program header information of the example: `readelf -l a.out`

```

2 Elf file type is DYN (Position-Independent Executable file)
Entry point 0x1100
4 There are 13 program headers, starting at offset 64

6 Program Headers:
  Type           Offset             VirtAddr           PhysAddr
8             FileSiz          MemSiz              Flags  Align
  PHDR           0x0000000000000040 0x0000000000000040 0x0000000000000040
10             0x00000000000002d8 0x00000000000002d8 R      0x8
  INTERP         0x0000000000000318 0x0000000000000318 0x0000000000000318
12             0x00000000000001c 0x00000000000001c R      0x1
    [Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]
14  LOAD          0x0000000000000000 0x0000000000000000 0x0000000000000000
             0x0000000000000760 0x0000000000000760 R      0x1000
16  LOAD          0x0000000000000100 0x0000000000000100 0x0000000000000100
             0x00000000000002cd 0x00000000000002cd R E    0x1000
18  LOAD          0x0000000000000200 0x0000000000000200 0x0000000000000200
             0x00000000000000ec 0x00000000000000ec R      0x1000
20  LOAD          0x00000000000002d9 0x00000000000003d9 0x00000000000003d9
             0x0000000000000280 0x0000000000000288 RW     0x1000
22  DYNAMIC       0x00000000000002da 0x00000000000003da 0x00000000000003da
             0x00000000000001f0 0x00000000000001f0 RW     0x8
24  NOTE          0x0000000000000338 0x0000000000000338 0x0000000000000338
             0x0000000000000030 0x0000000000000030 R      0x8
26  NOTE          0x0000000000000368 0x0000000000000368 0x0000000000000368
             0x0000000000000044 0x0000000000000044 R      0x4
28  GNU_PROPERTY  0x0000000000000338 0x0000000000000338 0x0000000000000338
             0x0000000000000030 0x0000000000000030 R      0x8
30  GNU_EH_FRAME  0x000000000000020c 0x000000000000020c 0x000000000000020c
             0x0000000000000034 0x0000000000000034 R      0x4
32  GNU_STACK     0x0000000000000000 0x0000000000000000 0x0000000000000000
             0x0000000000000000 0x0000000000000000 RW     0x10
34  GNU_RELRO     0x00000000000002d9 0x00000000000003d9 0x00000000000003d9
             0x0000000000000270 0x0000000000000270 R      0x1

36 Section to Segment mapping:
38 Segment Sections...
   00
40   01      .interp
   02      .interp .note.gnu.property .note.gnu.build-id .note.ABI-tag .
      gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .
      rela.plt
42   03      .init .plt .plt.got .plt.sec .text .fini
   04      .rodata .eh_frame_hdr .eh_frame
44   05      .init_array .fini_array .dynamic .got .data .bss
   06      .dynamic
46   07      .note.gnu.property
   08      .note.gnu.build-id .note.ABI-tag
48   09      .note.gnu.property

```

```

10      .eh_frame_hdr
50      11
      12      .init_array .fini_array .dynamic .got

```

programsection.txt

### 3 Linker

This is the final stage of compilation.

We know that library functions are not a part of any C program. Thus, the compiler doesn't know the operation of any function, whether it be `std::cin>>` or `std::cout <<`. The definitions of these functions are stored in their respective library but the compiler can link them. So, when we write `#include`, it includes `iostream` library which gives access to Standard Input and Output. The linker links the object files to the library functions and the program becomes an executable file. Here, `a.out` will be created in an executable format.

In this stage, the final executable ELF is generated. Linker uses all object files created for each programming file in earlier stages as inputs.

It takes a linker descriptor file as input. The linker combines sections from all object files into one final ELF. It instructs the linker descriptor file as to what addresses are to be provided for each section.

It also resolves the symbols used by each object file. If any symbol is not defined, then linker gives an error. These are the stages by which a final executable is generated.

### 4 Program execution

If we execute any C program, then `main()` is always the entry point. This is because a startup code (provided by compiler) is always linked to the original program by linker. This startup code consists of a startup function, which usually is `__start` or `Reset_handler`. This is placed at the address that the CPU will execute at boot. This function calls the `main()` function, which is the starting point of the application.

Whenever we give the command to execute a program, the loader will load the ELF file into memory and inform the CPU with the starting point of the address where this program is loaded.

### 5 Code Listings

```

1 #include <iostream>
  using namespace std;
3 int main()
  {
5     std::cout << "Hello World!" << std::endl;
  }

```

hello-world.cpp