



Graph Algorithms 2

Outline


- Prim's algorithm
- Kruskal's algorithm
- A* Search algorithm
- Max Flow Problem
- Bipartite graph

Tree vs Graph

- Trees are undirected graphs (not all undirected graphs are trees)
- Trees do not have cycles
- Trees are connected graphs, connected acyclic undirected graphs
- If a tree has n vertices, then it has $n - 1$ edges
 - Less than $n - 1$ edges \rightarrow Disconnected
 - More than $n - 1$ edges \rightarrow There is a cycle






Not a tree




Not a tree

Spanning Trees


- Input: a **connected** undirected graph G with n vertices
- Find such a subgraph:
 - Keep all the n vertices
 - Keep $n - 1$ edges
 - The subgraph is connected
- The subgraph is a spanning tree
 - Not unique
 - For G with positive edge weights: **Minimum spanning tree (MST)** is a spanning tree that minimizes the sum of weights




Sum of weights:16



Sum of weights:26





Sum of weights:25



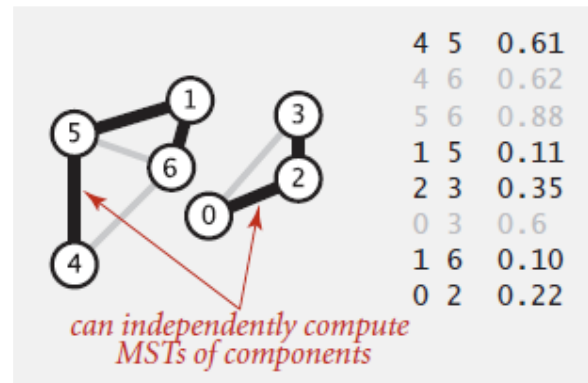
Sum of weights:21

Minimum spanning tree (MST)

- What if edge weights are not all distinct?
 - Greedy MST algorithms still correct if equal weights are present!




- What if graph is not connected?
 - No MST, but there is a minimum spanning forest = MST of each component



- How to represent the MST
 - A list of edges (with weights)

Prim's Algorithm


- Basic idea: Grow the tree in successive stages
- Initially, the tree has one vertex and no edge
- In each iteration, add one vertex and one edge to the tree
- Throughout, maintain the properties of trees:
 - Connectivity
 - No cycle: disregard if the vertices are visited
- The algorithm runs in n iterations (n is the number of vertices)



\mathcal{U} : vertices of spanning tree

Iteration 1


- Pick any vertex in the graph
- Maybe pick v_1
- Add v_1 to \mathcal{U}



$$\mathcal{U} = \{v_1\}$$

Iteration 2


- The edges connecting \mathcal{U} to $\mathcal{V} \setminus \mathcal{U}$:
 $e_{1,2}, e_{1,3}, e_{1,4}$
- Among them, $e_{1,4}$ has the smallest weight, record it
- Add v_4 to \mathcal{U}



$$\mathcal{U} = \{v_1\}$$

Iteration 3

- The edges connecting \mathcal{U} to $\mathcal{V} \setminus \mathcal{U}$:
 $e_{1,2}, e_{1,3}$
 $e_{4,2}, e_{4,3}, e_{4,5}, e_{4,6}, e_{4,7}$
- Among them, $e_{1,2}$ has the smallest weight, record it
- Add v_2 to \mathcal{U}



$$\mathcal{U} = \{v_1, v_4, v_2\}$$


Iteration 4

- The edges connecting \mathcal{U} to $\mathcal{V} \setminus \mathcal{U}$:

$e_{1,3}$
 $e_{4,3}, e_{4,5}, e_{4,6}, e_{4,7}$

$e_{2,5}$


- Among them, $e_{4,3}$ has the smallest weight, record it
- Add v_3 to \mathcal{U} .



$$\mathcal{U} = \{v_1, v_4, v_2, v_3\}$$

Iteration 5

- The edges connecting \mathcal{U} to $\mathcal{V} \setminus \mathcal{U}$:
 $e_{4,5}, e_{4,6}, e_{4,7}$
 $e_{2,5}$
 $e_{3,6}$
- Among them, $e_{4,7}$ has the smallest weight, record it
- Add v_7 to \mathcal{U}



$$\mathcal{U} = \{v_1, v_4, v_2, v_3, v_7\}$$

Iteration 6

- The edges connecting \mathcal{U} to $\mathcal{V} \setminus \mathcal{U}$:


$e_{4,5}, e_{4,6}$

$e_{2,5}$

$e_{3,6}$

$e_{7,5}, e_{7,6}$


- Among them, $e_{7,6}$ has the smallest weight, record it
- Add v_6 to \mathcal{U}



$$\mathcal{U} = \{v_1, v_4, v_2, v_3, v_7, v_6\}$$

Iteration 7

- The edges connecting \mathcal{U} to $\mathcal{V} \setminus \mathcal{U}$:
 $e_{4,5}$
 $e_{2,5}$
 $e_{7,5}$
- Among them, $e_{7,5}$ has the smallest weight, record it
- Add v_5 to \mathcal{U}
- After this
 - Now $\mathcal{U} = \mathcal{V}$. (All the vertices have been added to \mathcal{U} .)
 - Return the tree:
 $\{e_{1,4}, e_{1,2}, e_{4,3}, e_{4,7}, e_{7,6}, e_{7,5}\}$



$$\mathcal{U} = \{v_1, v_4, v_2, v_3, v_7, v_6, v_5\}$$

Prim's Algorithm

1. Initialize

- Let \mathcal{T} (the set of edges in the MST) be an empty set
- Create an array `minWeight` of size n (number of nodes) to store the minimum weight of an edge that connects each node to the MST
- Initialize `minWeight[start] = 0` (starting node with 0 cost to itself) and all other values in `minWeight` to infinity
- Create an array `inMST` of size n to keep track of whether each node is in the MST, initialize all values in `visited` to false

2. While \mathcal{T} has fewer than $n-1$ edges

- Find the unvisited node u with the minimum value in `minWeight`
- Mark u as visited and add it to the MST
- For each neighbor v of u :
 - If v is unvisited and the edge weight (u,v) is less than `minWeight[v]`
update `minWeight[v]` to the weight of (u,v)

3. Return \mathcal{T}

```
// Prim's algorithm without priority queue
void prim(int start, vector<vector<int>>& graph) {
    int n = graph.size();
    vector<bool> inMST(n, false);    // Track nodes in MST
    vector<pair<int, int>> mstEdges; // Store edges in MST
    int totalCost = 0;
    inMST[start] = true;    // Add starting node to MST
    for (int count = 1; count < n; ++count) {    // Repeat until all nodes are in MST
        int minWeight = INT_MAX;
        int u = -1, v = -1;
        // Find the minimum weight edge (u, v) with u in MST and v not in MST
        for (int i = 0; i < n; ++i) {
            if (inMST[i]) {
                for (int j = 0; j < n; ++j) {
                    if (!inMST[j] && graph[i][j] < minWeight) {
                        minWeight = graph[i][j];
                        u = i;
                        v = j;
                    }
                }
            }
        }
        // Add edge (u, v) to MST
        if (u != -1 && v != -1) {
            inMST[v] = true;
            totalCost += minWeight;
            mstEdges.push_back({u, v});
        }
    }
}
```

Improving using priority queue

- Challenge: Find the min weight edge with exactly one endpoint in \mathcal{U}
- Solution: Maintain a PQ of edges with (at least) one endpoint in \mathcal{U}
 - Key = edge; priority = weight of edge
 - Delete-min to determine next edge $e_{u,v}$ to add to \mathcal{U}
 - Disregard if both endpoints u and v are unvisited (both in \mathcal{U})
 - Otherwise, let v be the unvisited vertex (not in \mathcal{U}):
 - add to PQ any edge incident to v (assuming other endpoint not in \mathcal{U})
 - add e to \mathcal{U} and mark v as visited

Minimum edge weight data structure	Time complexity
adjacency matrix, searching	$O(n^2)$
binary heap, priority queue	$O(m \log(n))$

n : # of vertices, m : # of edges


```
MST-Prim(G, V, E){
```

```
  for(each x∈V){
```

```
    cost(x)=∞;
```

```
    parent(x)=Null;
```

```
  }
```

```
  Choose a node u to be the source or starting point;
```

```
  cost(u)=0;
```

```
  Insert all vertices to a priority queue PQ;
```

```
  while (PQ≠∅){
```

```
    x = PQ.Extract_Min();
```

```
    for (each neighbor, y of x){
```

```
      if (y!∈U && w(x,y)<cost(y)){
```

```
        cost(y)=w(x,y);
```

```
        parent(y)=x;
```

```
        PQ.Decrease_Key(y, cost(y));
```

```
      }
```

```
    }
```


```
  }}
```

Kruskal's algorithm

- Prim's Algorithm: vertex-wise; Kruskal's algorithm: edge-wise
 - Basic idea: Maintain a forest, i.e., a collection of trees
 - Initially, there are n trees; every vertex is a tree
 - Each iteration examines one edge; the edge may be chosen so that two trees are merged
 - Stop when there is only one tree
 - The algorithm runs in at most m iterations (Because there are m edges)
-
-

Preparation

- Build a queue of edges
- Sorted so that the weights are in ascending order




$\mathcal{T} = \emptyset$. (Record the selected edges.)

Edge	Weight
(1, 4)	1
(6, 7)	1
(1, 2)	2
(3, 4)	2
(2, 4)	3
(1, 3)	4
(4, 7)	4
(3, 6)	5
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

Iteration 1

- Dequeue and get the edge (1, 4)
- v_1 and v_4 are not in the same tree.
- Thus **accept** edge (1, 4)
- Append (1, 4) to \mathcal{T}




$$\mathcal{T} = \{e_{1,4}\}$$

Edge	Weight
(1, 4)	1
(6, 7)	1
(1, 2)	2
(3, 4)	2
(2, 4)	3
(1, 3)	4
(4, 7)	4
(3, 6)	5
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

Iteration 2

- Dequeue and get the edge (6, 7)
- v_6 and v_7 are not in the same tree
- Thus **accept** edge (6, 7)
- Append (6, 7) to \mathcal{T}




$$\mathcal{T} = \{e_{1,4}, e_{6,7}\}$$

Edge	Weight
(6, 7)	1
(1, 2)	2
(3, 4)	2
(2, 4)	3
(1, 3)	4
(4, 7)	4
(3, 6)	5
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

Iteration 3

- Dequeue and get the edge (1, 2)
- v_1 and v_2 are not in the same tree
- Thus **accept** edge (1, 2)
- Append (1, 2) to \mathcal{T}




$$\mathcal{T} = \{e_{1,4}, e_{6,7}, e_{1,2}\}$$

Edge	Weight
(1, 2)	2
(3, 4)	2
(2, 4)	3
(1, 3)	4
(4, 7)	4
(3, 6)	5
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

Iteration 4

- Dequeue and get the edge (3, 4)
- v_3 and v_4 are not in the same tree
- Thus **accept** edge (3, 4)
- Append (3, 4) to \mathcal{T}




$$\mathcal{T} = \{e_{1,4}, e_{6,7}, e_{1,2}, e_{3,4}\}$$

Edge	Weight
(3, 4)	2
(2, 4)	3
(1, 3)	4
(4, 7)	4
(3, 6)	5
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

Iteration 5

- Dequeue and get the edge (2, 4)
- v_2 and v_4 are in the same tree
- Thus **reject** edge (2, 4)




$$\mathcal{T} = \{e_{1,4}, e_{6,7}, e_{1,2}, e_{3,4}\}$$

Edge	Weight
(2, 4)	3
(1, 3)	4
(4, 7)	4
(3, 6)	5
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

Iteration 6

- Dequeue and get the edge (1, 3)
- v_1 and v_3 are in the same tree
- Thus **reject** edge (1, 3)




$$\mathcal{T} = \{e_{1,4}, e_{6,7}, e_{1,2}, e_{3,4}\}$$

Edge	Weight
(1, 3)	4
(4, 7)	4
(3, 6)	5
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

Iteration 7

- Dequeue and get the edge (4, 7)
- v_4 and v_7 are not in the same tree
- Thus accept edge (4, 7)
- Append (4, 7) to \mathcal{T}




$$\mathcal{T} = \{e_{1,4}, e_{6,7}, e_{1,2}, e_{3,4}, e_{4,7}\}$$

Edge	Weight
(4, 7)	4
(3, 6)	5
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

Iteration 8

- Dequeue and get the edge (3, 6)
- v_3 and v_6 are in the same tree
- Thus **reject** edge (3, 6)




Edge	Weight
(3, 6)	5
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

$$\mathcal{T} = \{e_{1,4}, e_{6,7}, e_{1,2}, e_{3,4}, e_{4,7}\}$$

Iteration 9

- Dequeue and get the edge (5, 7)
- v_5 and v_7 are not in the same tree
- Thus **accept** edge (5, 7)
- Append (5, 7) to \mathcal{T}
- After this
 - All the vertices are connected
 - Return the edges \mathcal{T}



Edge	Weight
(5, 7)	6
(4, 5)	7
(4, 6)	8
(2, 5)	10

$$\mathcal{T} = \{e_{1,4}, e_{6,7}, e_{1,2}, e_{3,4}, e_{4,7}, e_{5,7}\}$$


Kruskal's Algorithm

1. Put all the edges of the input graph into a queue
 2. Sort the queue so that the weights are in ascending order
 3. Let set \mathcal{T} (which stores the selected edges) be the empty set
 4. While \mathcal{T} has fewer than $n-1$ edges
 - Get an edge: $e_{u,v} \leftarrow \text{dequeue}()$
 - If u and v are in different trees, then add $e_{u,v}$ to \mathcal{T} and merge the two trees.
 5. Return \mathcal{T}
-


Using Disjoint Sets Data Structure

- How to decide whether two vertices are in the same tree?
- Solution: Using disjoint sets data structure. Put vertices of a tree in the same set. Deciding whether two vertices belong to the same set costs near $O(1)$ time. **Find()**
- How to merge two trees? **Union()**

Find()




Union()




Disjoint Subsets

- Definition: $A \cap B = \emptyset \rightarrow A$ and B are disjoint subsets.
- Example: $A = \{1, 2, 3\}$ and $B = \{4, 5\}$; share no same item \rightarrow disjoint.
- Object set: $\emptyset, 1, 2, 3, 4, 5$
- Disjoint subset: $\{\emptyset\} \{1\} \{2\} \{3\} \{4\} \{5\} \{6\}$
 - Set A: $\{\emptyset\}$
 - Set B: $\{1\}$
 - Set C: $\{2, 3, 4\}$
 - Set D: $\{5, 6\}$
- Common operations:
 - Find: $\text{find}(\text{item}) \rightarrow$ Set ID, i.e., $\text{find}(3) \rightarrow C$.
 - Union: $\text{union}(A, B) \rightarrow$ new subset $A \cup B = \{\emptyset, 1\}$.
- How to support fast find and union?
- Idea: use a tree to represent each subset.




Disjoint Subsets – Trees

- Disjoint subset: $\{0\}$ $\{1\}$ $\{2\}$ $\{3\}$ $\{9\}$ $\{5\}$ $\{6\}$
- Common operations:
 - Find: $\text{find}(\text{item}) \rightarrow \text{Set ID}$, i.e., $\text{find}(3) \rightarrow C$.
 - Union: $\text{union}(A, B) \rightarrow \text{new subset } A \cup B = \{0, 1\}$.
- Tree id: the root item.
 - Each subset has one item as id. Unique?
- Data struct for each item: [Item id, Parent item id]
 - Root item: parent item is itself.
- Set 0: $[0, 0] \rightarrow$ only one item in the subset.
- Set 1: $[1, 1] \rightarrow$ only one item in the subset.
- Set 2: a) $[9, 3]$; b) $[3, 2]$; c) $[2, 2]$. Who is the root?
- Set 3: a) $[6, 5]$; b) $[5, 5]$.




Disjoint Subsets – Find and Union

- Many ways for the find and union operations.
- Let us start with some intuitive ones.
- With the trees, what is the time complexity of find(\cdot)?
 - Worst case, $O(V)$ items in a tree with one node per level.
 - Finding the leaf item, so take $O(V)$ checks to the root.
- What is the time complexity of union(\cdot, \cdot).
 - Relink both trees. $O(1)$.
 - i.e., union(Set 0, Set 1).
 - Item 1's pointer to item 0.
- Fine for union(), but not for find().
- How to make find() faster?
 - A lot options. We cover a few here.




Disjoint Subsets – Faster Find

- Data struct for each item: [Item id, Parent item id, **root id**].
 - Allocate some space to record the root id for each item in the tree.
 - Object set: 0, 1, 2, 3, 5, 6, 9
 - Root id: 0, 1, 2, 2, 5, 5, 2
- Here, find() can be in $O(1)$ time.
- But how about union? Still $O(1)$?
 - Merge two trees together with relinking.
 - Update all root id. $O(V)$ operation for large trees.
- Now we presented two options:
 - Quick union $O(1)$ with slow find $O(V)$.
 - Quick find $O(1)$ with slow union $O(V)$.
- Question: why do we have a tree looks like a list?
 - For list of length $O(V)$, $O(V)$ time traverse.
 - Binary tree, can be $O(\log V)$ levels, much faster. But how?



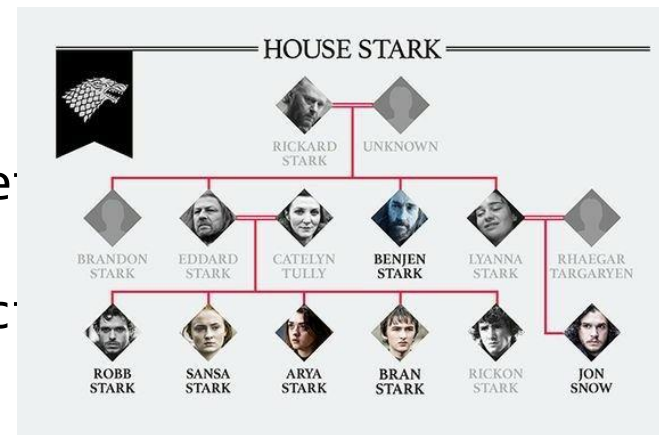

Disjoint Subsets – Maintain a Balanced Tree

- Initially, **build a flat tree**.
- Throughout the further unions, **keep flat**, or balanced.
- Check our example to merge $\{5, 6\}$ and $\{2, 3, 9\}$.
 - One option gives us flat tree.
 - The other option makes the tree very imbalanced.
- Intuitively, root shall come from the larger subset.
 - i.e., 2 shall be the root compared to 5.
- For implementation, how to do this?
 - How can we say which tree is larger or smaller?
- Choice 1: record tree size in root.
 - Tree size: the number of items in the tree.
 - Root from the tree with more items.
 - Smaller tree becomes a subtree.
 - Update tree size, i.e., $\text{int } A + \text{int } B$.
 - Time for `find()` being pushed to $O(\log|V|)$.




Disjoint Subsets – Applications

- Choice 2: record tree depth, or rank, in root.
 - Shorter tree becomes a subtree.
 - Key idea is still to balance the tree.
 - Time for find() being pushed to $O(\log|V|)$.
- More advanced options available.
 - i.e., path compression, path halving, path splitting.
 - Different options, or the combination of options, give different complexity.
 - Read references if interested.
- Applications:
 - Network connectivity.
 - Local networks -> global ne
 - Image segmentation.
 - From one pixel to one objec
 - Least common ancestor.



A* Search algorithm

- DFS, BFS, Dijkstra...
 - **Expand “blindly”** by exploring nodes without any “intelligence” guidance towards the goal
 - **Consider all directions as equally likely** to lead to the goal or follow a **fixed order** of exploration.
 - Do not have a "sense" of direction toward the goal, meaning they operate without any preference for paths that might be closer to the target
 - Time consuming if the graph is big even with Bidirectional or Iterative Deepening improvements
- Other Improvements
 - Heuristic search: rank the directions/options based on importance (priority)
 - Only expand limited options, consider more options if cannot find – Iterative Widening
 - Pruning, stop expanding for options that unlikely/unfeasible to reach the goal



Heuristic Search

- Key: Heuristic Function
- The heuristic function assigns a priority to each direction or option based on the current state, guiding the search towards the goal.
- A more accurate heuristic function leads to faster convergence on the optimal solution by reducing unnecessary exploration.
- No one-size-fits-all heuristic; need to tailored to the specific problem.
- The heuristic should be computationally efficient, as a costly heuristic can negatively impact search performance, e.g. Manhattan distance,

$$h(s) = |x_s - x_g| + |y_s - y_g|$$

```

bfs(s) {
    q = new queue()
    q.push(s), visited[s] = true
    while (!q.empty()) {
        u = q.pop() - check if meet the goal
        for each edge(u, v) {
            if (!visited[v]) {
                q.push(v)
                visited[v] = true
            }
        }
    }
}

```

```

greedyBestFirstSearch(s) {
    q = new priority_queue()
    q.push(s, h(s)), visited[s] = true
    while (!q.empty()) {
        u = q.pop() - check if meet the goal
        for each edge(u, v) {
            if (!visited[v]) {
                q.push(v, h(s))
                visited[v] = true
            }
        }
    }
}

```

- Cost Function:

- $g(x)$: actual moving cost from the start node to the current node x
- $h(x)$: heuristic function, representing the estimated cost from the current node x to the goal
- The **priority queue** returns the node x with the minimum $f(x)$, prioritizing nodes that appear closer to the goal based on $f(x)$

```

function AStar(Graph, start, goal):
    create vertex priority queue Q
    distTo[start] ← 0 // g(s) = 0 for the start node
    Q.add_with_priority(start, h(start)) // priority = h(start)
    for each vertex v in Graph.Vertices:
        if v ≠ start
            prev[v] ← UNDEFINED
            distTo[v] ← INFINITY // Set g(v) = ∞ initially
            Q.add_with_priority(v, INFINITY)

    while Q is not empty:
        u ← Q.extract_min() // Node with lowest f(u) = g(u) + h(u)
        if u == goal: // Goal reached
            return distTo, edgeTo // Shortest path found

        for each neighbor v of u:
            tentative_gScore ← distTo[u] + Graph.Edges(u, v)
            if tentative_gScore < distTo[v]:
                edgeTo[v] ← u // Track the path
                distTo[v] ← tentative_gScore // Update g(v) = g(u) + cost(u, v)
                fScore ← distTo[v] + h(v) // Calculate f(v) = g(v) + h(v)
                Q.decrease_priority(v, fScore) // Update priority in Q with f(v)

    return 0 // If Q is empty and goal not found


```

IDA* (Iterative Deepening A*):

- **Challenge with A*:** While effective, A* can consume large amounts of memory, storing all expanded nodes in the priority queue and open/closed lists
 - **Solution with IDA*:** Similar to Iterative Deepening Depth-First Search (IDDFS), IDA* applies a **cost limit instead of a depth limit**. Each iteration explores paths within a given $f(x)$ threshold (cost limit), gradually increasing the limit in subsequent iterations
 - **Advantage:** IDA* significantly reduces memory usage compared to A*, as it doesn't need to store the entire search tree, making it suitable for memory-constrained environments
-


Max Flow Problem

- Send water from the source s to the sink t
- The edges are pipes which have certain capacities, e.g., $4m^3/s$
- How much water can flow from source s to the sink t at most?
- Inputs: A weighted directed graph, the source s , and the sink t
- Goal: Send as much water as possible from s to t
- Constraints:
 - Each edge has a weight (i.e., the capacity of the pipe)
 - The flow must not exceed the capacity




Initialization

- Augmenting path: a path from s to t that does not contain cycles



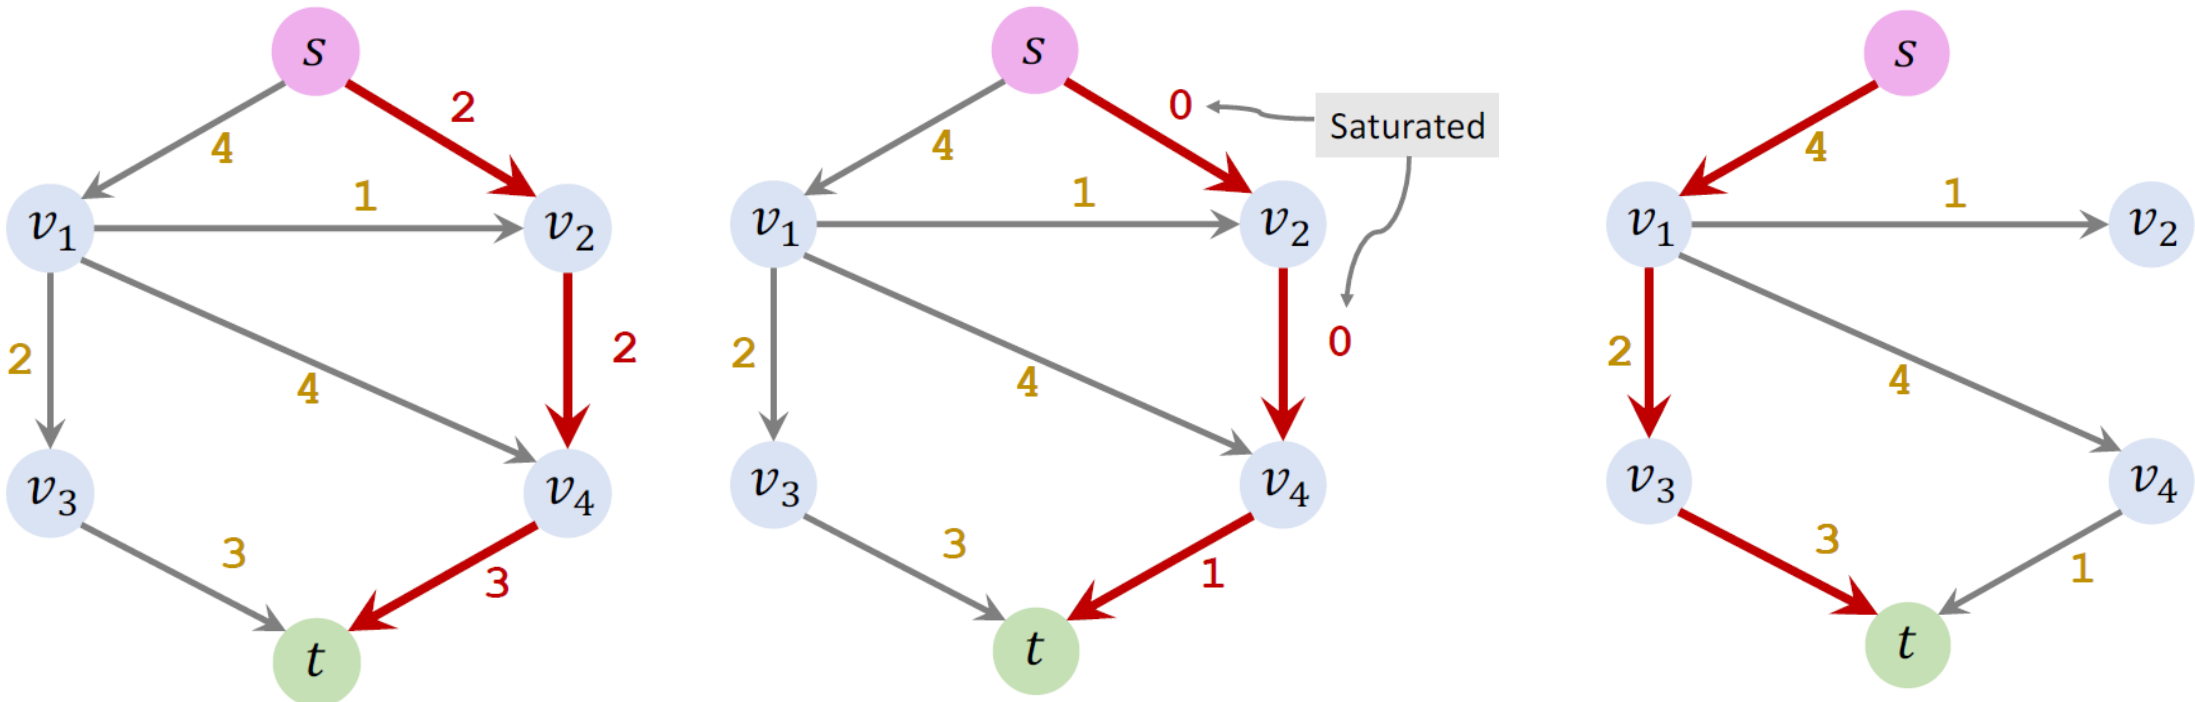
Original Graph



Residual Graph

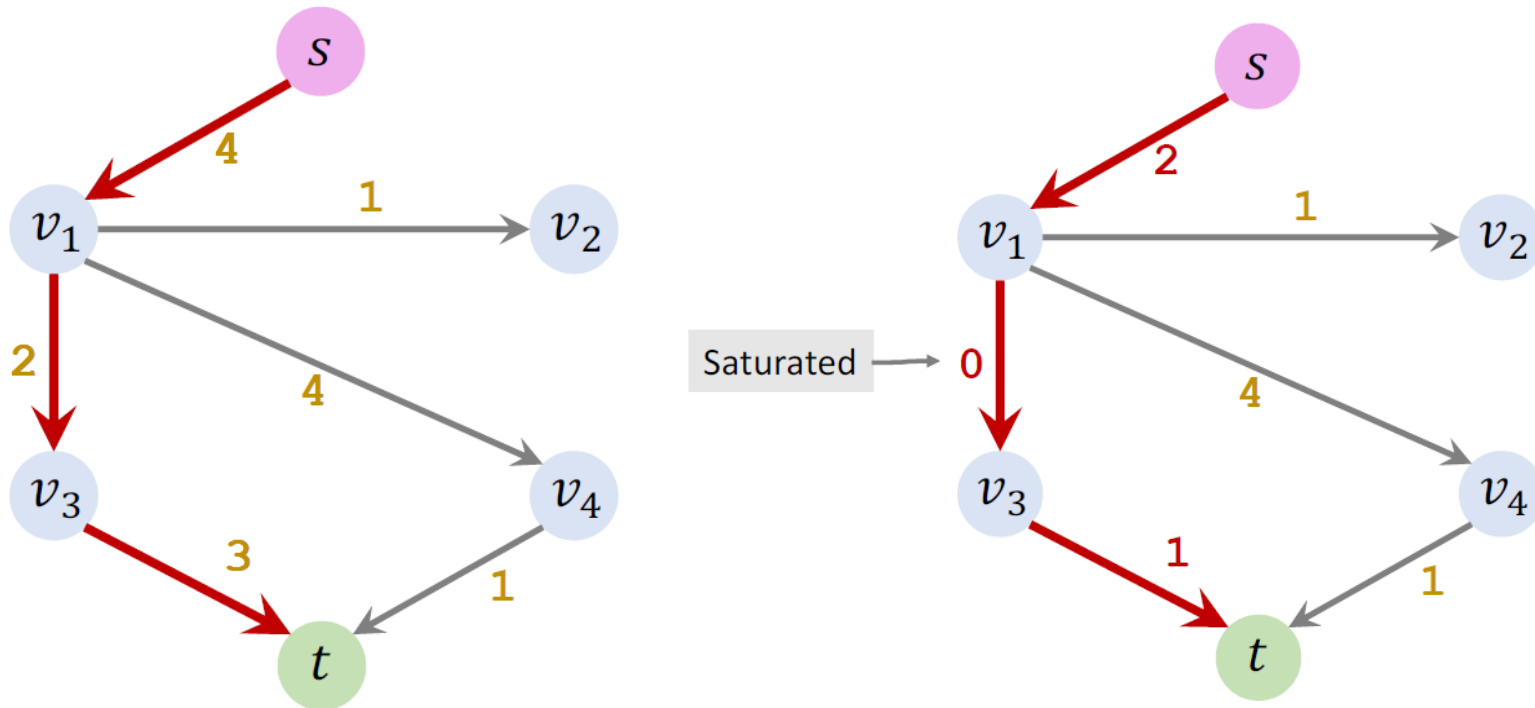
Iteration 1: find an augmenting path and update residuals

- Working on the **residual graph**
- Augmenting path: a path from s to t that does not contain cycles
- Found path $s \rightarrow v_2 \rightarrow v_4 \rightarrow t$ (Bottleneck capacity = 2)
- Update residuals: -2 for the edges, remove saturated edges






Iteration 2: find an augmenting path and update residuals

- Working on the **residual graph**
- Augmenting path: a path from s to t that does not contain cycles
- Found path $s \rightarrow v_1 \rightarrow v_3 \rightarrow t$ (Bottleneck capacity = 2)
- Update residuals: -2 for the edges, remove saturated edges



Iteration 2: find an augmenting path and update residuals



- Working on the **residual graph**
- Augmenting path: a path from s to t that does not contain cycles
- Found path $s \rightarrow v_1 \rightarrow v_4 \rightarrow t$ (Bottleneck capacity = 1)
- Update residuals: -1 for the edges, remove saturated edges



Cannot find any path from source to sink

Flow = Capacity - Residual

- Amount of flow: 5, outward flow of s == inward flow of t == 5




$$\text{Flow} = \text{Capacity} - \text{Residual}.$$

Simple solution


1. Build a residual graph; initialize the residuals to the capacity
2. While augmenting path can be found:
 1. Find an augmenting path (on the residual graph)
 2. Find the bottleneck capacity x in the augmenting path
 3. Update the residuals ($\text{residual} \leftarrow \text{residual} - x$)
3. **Flow = Capacity - Residual**

This simple solution may fail


- Always finds the blocking flow, may not be the maximum flow
- A flow is blocking flow if no more flow from source to sink can be found
- The “pipes” are blocked
- Maximum flow is also blocking flow
- Once a bad path is selected, the simple solution cannot make corrections



Maximum Flow




Not Maximum Flow




Amount of flow: 4

Ford-Fulkerson Algorithm

- Key idea: allow correction by adding backward path






Original Graph



Residual Graph


Iteration 1

- Working on the **residual graph**
- Found augmenting path $s \rightarrow v_1 \rightarrow v_4 \rightarrow t$ (Bottleneck capacity = 3)
- Update residuals: -3 for the edges, remove saturated edges
- **Add a backward path** $t \rightarrow v_4 \rightarrow v_1 \rightarrow s$ with weight= 3




Iteration 2

- Working on the **residual graph**
- Found augmenting path $s \rightarrow v_1 \rightarrow v_3 \rightarrow t$ (Bottleneck capacity = 1)
- Update residuals: -1 for the edges, remove saturated edges
- **Add a backward path** $t \rightarrow v_3 \rightarrow v_1 \rightarrow s$ with weight= 1




Iteration 3

- Working on the **residual graph**
- Found augmenting path $s \rightarrow v_2 \rightarrow v_4 \rightarrow v_1 \rightarrow v_3 \rightarrow t$ (Bottleneck capacity = 1)
- Update residuals: -1 for the edges, remove saturated edges
- **Add a backward path** $t \rightarrow v_3 \rightarrow v_1 \rightarrow v_4 \rightarrow v_2 \rightarrow s$ with weight= 1
- Cannot find any path anymore





Flow = Capacity - Residual

- Max flow = 5




Original Graph




Residual Graph

Ford-Fulkerson bad case


- The amount of the max flow is 200
- Ford-Fulkerson algorithm may take 200 iterations to find the max flow
- In every iteration, the amount of flow increases by 1




Original Graph



Residual Graph



Iteration 1



Iteration 2

Time Complexity


- Each iteration increases the amount of flow by at least 1
- Thus, # Iterations \leq Amount of Max Flow
- It takes $O(m)$ time to find a path in unweighted graph (Ignore the weights in the residual graph) (m : # of edges, f : amount of max flow)
- Thus, the per-iteration time complexity is $O(m)$
- The worst-case time complexity is $O(fm)$

Edmonds-Karp Algorithm


- Ford-Fulkerson:
 1. Build a residual graph; initialize the residuals to the capacities
 2. While augmenting path can be found:
 1. Find an augmenting path (on the residual graph)
Improvement: Find the shortest augmenting path (on the residual graph)
 2. Find the bottleneck capacity x on the augmenting path
 3. Update the residuals ($\text{residual} \leftarrow \text{residual} - x$)
 4. Add a backward path (Along the path, all edges have weights of x)
- Improvement is Edmonds-Karp
 - When finding path, regard the residual graph as unweighted
 - This can be found by a BFS, where we apply a weight of 1 to each edge
 - Time complexity: $O(m^2n)$ (m is #edges; n is #vertices)

Dinic's Algorithm


- Based on **level graph**: only edges connect to the next BFS level. Assign levels to all nodes, level of a node is shortest distance (in terms of number of edges) of the node from source.




Original Graph



Level Graph




Original Graph




Level Graph

Dinic's Algorithm


- Construct level graph
- Find blocking flow in level graph
 - Blocking flow: if no more flow from source to sink can be found
 - Blocking flow can be found using the simple solution



Original Graph





Residual Graph



Level Graph


Iteration 1

- Construct level graph
- Find blocking flow in level graph
- Update the residual graph, remove saturated edges
- Add flows to the residual graph as backward paths




Iteration 2


- Construct level graph
- Find blocking flow in level graph
- Update the residual graph, remove saturated edges
- Add flows to the residual graph as backward paths



Old Residual Graph



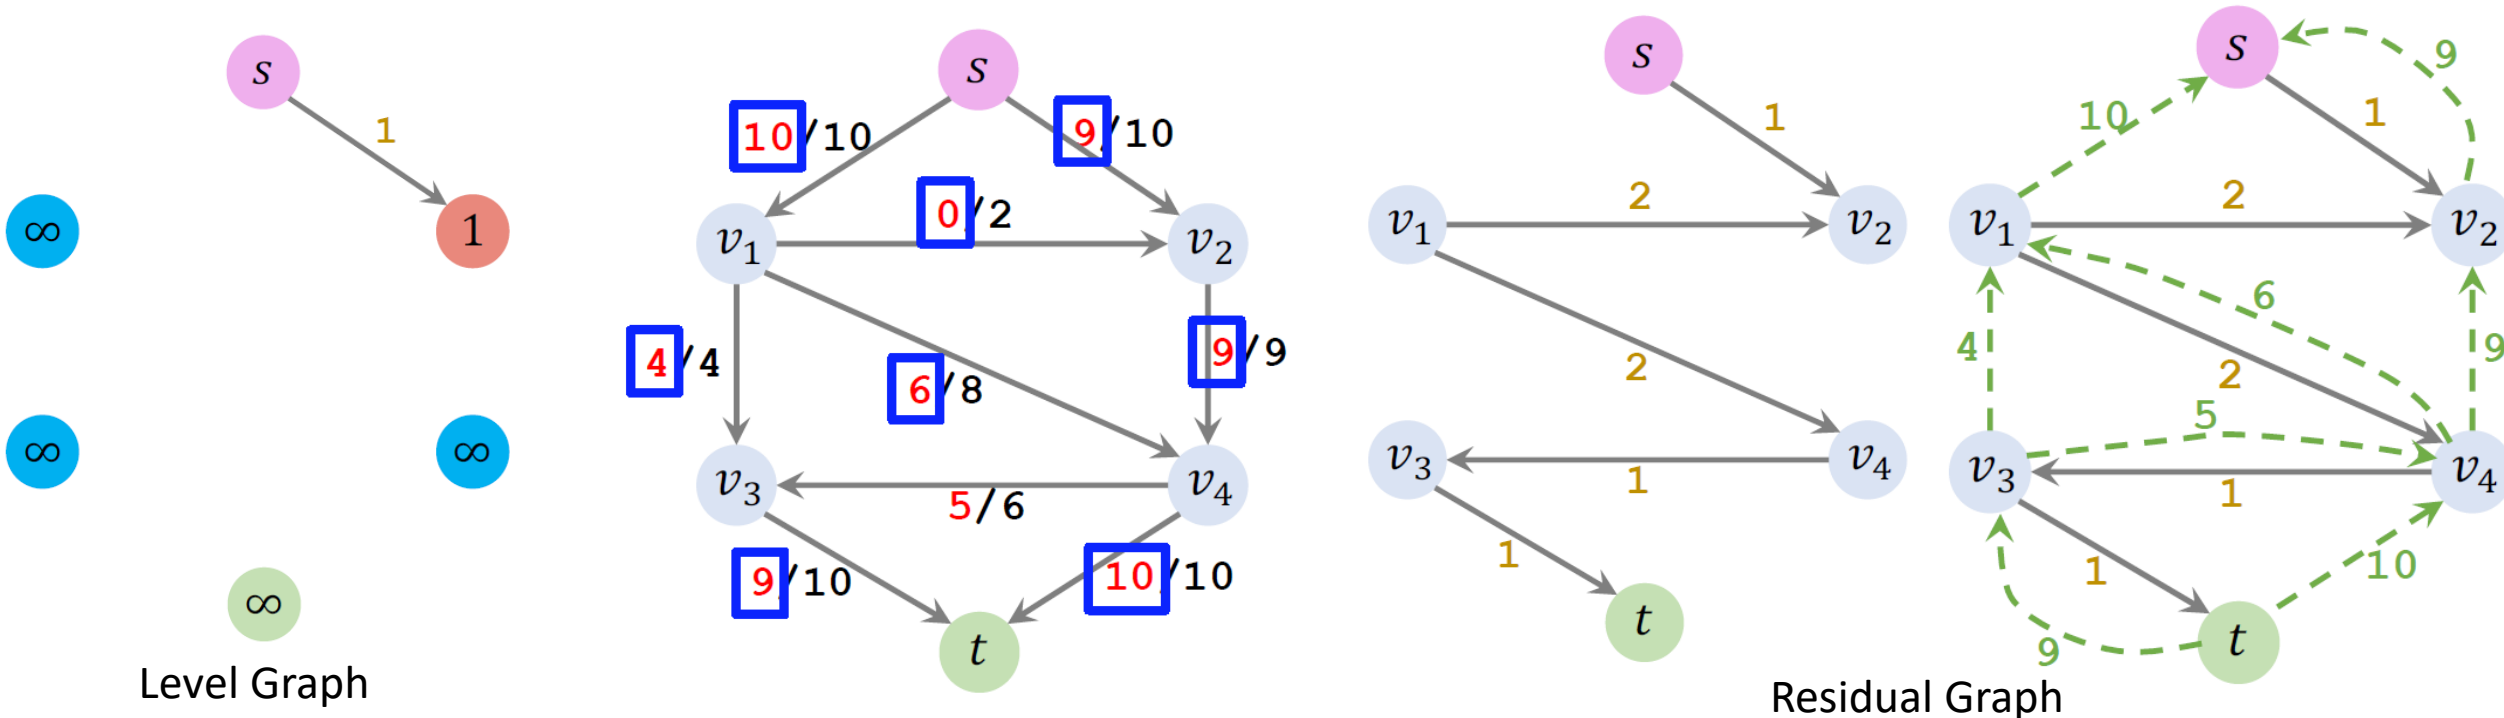
Level Graph



Residual Graph

Iteration 3

- Construct level graph
- Find blocking flow in level graph
- On the level graph, no flow can be found
- Flow = Capacity - Residual




Dinic's algorithm: Time Complexity

1. Initially, the residual graph is a copy of the original graph
 2. Repeat:
 1. Construct the level graph of the residual graph
 2. Find a blocking flow on the level graph
 3. Update the residual graph (update the weights, remove saturated edges, and add backward edges)
- Time complexity: $O(mn^2)$, (n is #vertices, m is #edges)
 - Dinic's algorithm has at most $n - 1$ iterations.
 - Per-iteration time complexity is $O(mn)$

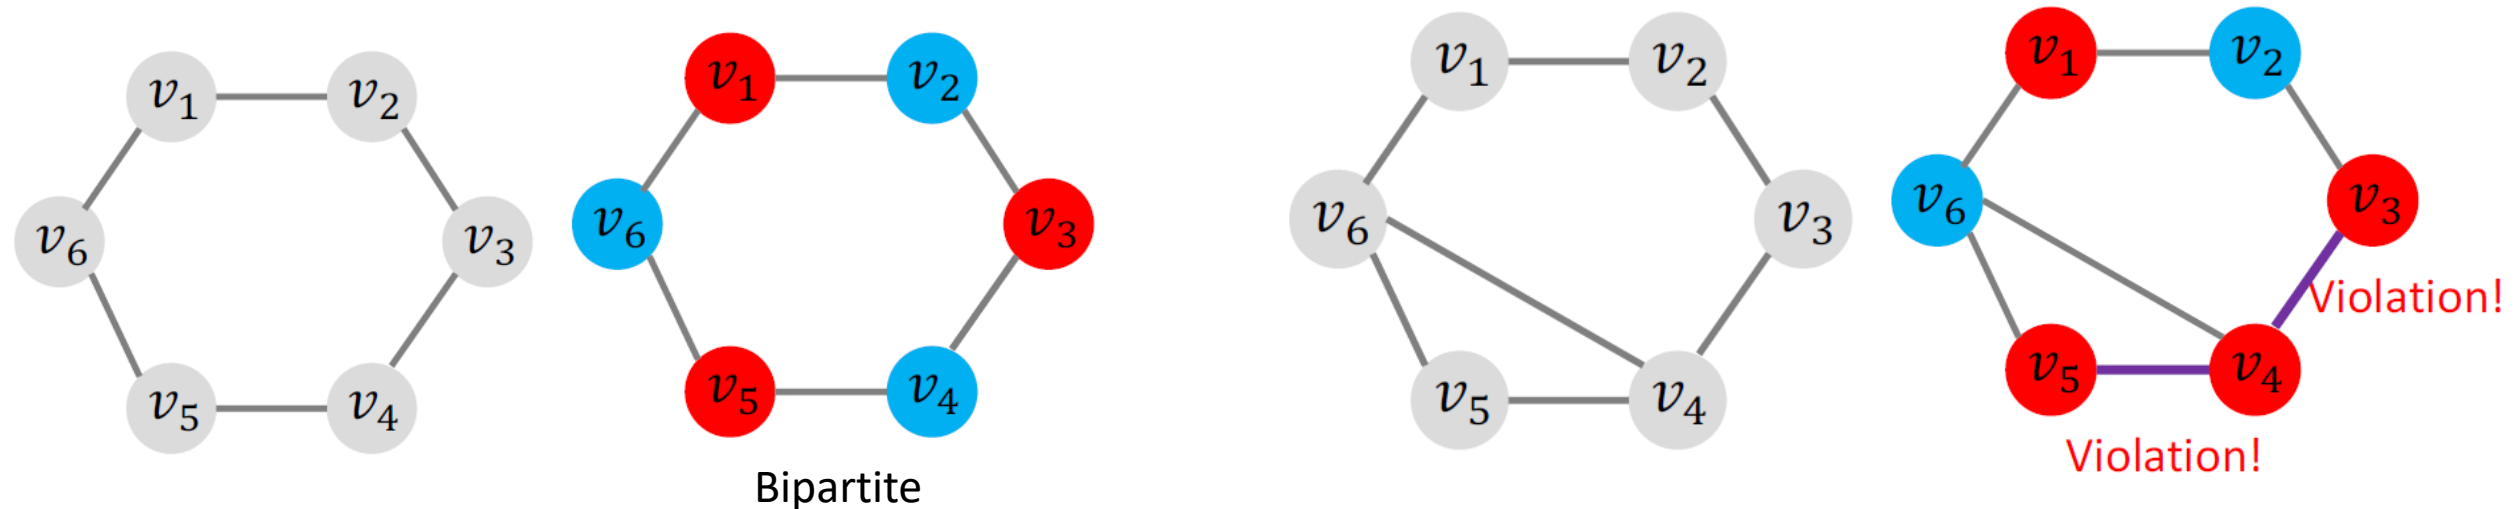
Bipartite Graph

- Bipartite graph: $\mathcal{G} = \mathcal{U}, \mathcal{V}, \mathcal{E}$
- All the edges are between \mathcal{U} and \mathcal{V}
- No edge between two vertices in \mathcal{U}
- No edge between two vertices in \mathcal{V}



Is the graph bipartite?

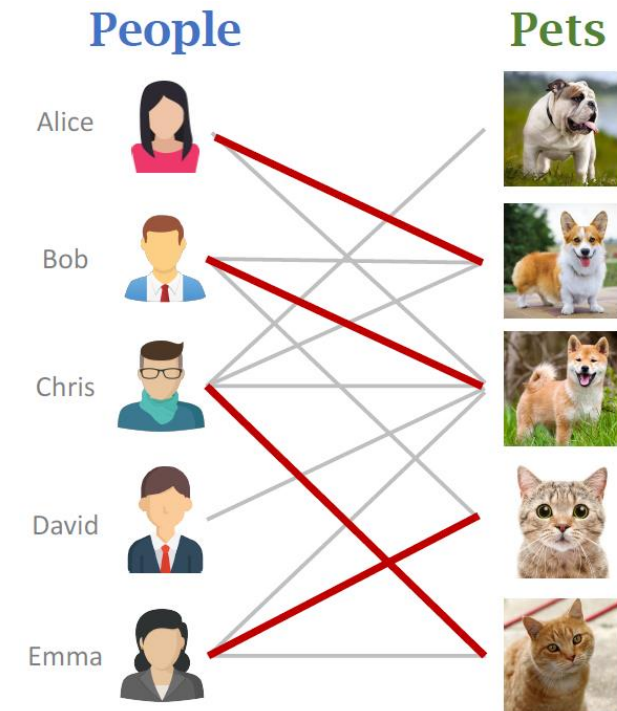
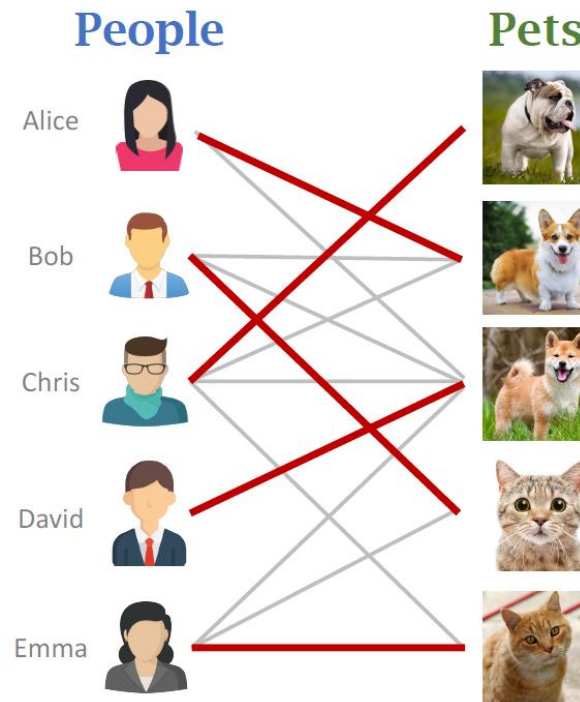
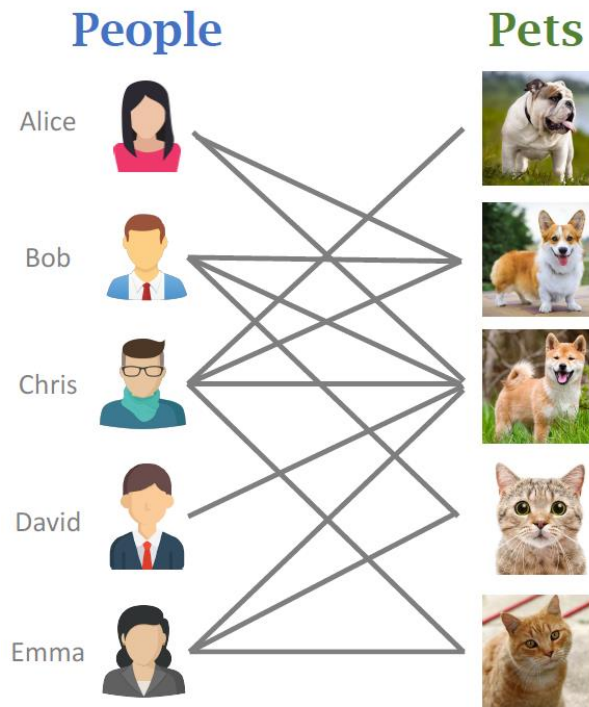
1. Select an arbitrary vertex and assign red color to it.
2. BFS to color neighbors, until all vertices are colored:
 1. Color red vertices' neighbors as blue
 2. Color blue vertices' neighbors as red
3. During the process, if a vertex has the same color as its neighbor, then output FALSE
3. If no violation is found, return TRUE in the end



Maximum-Cardinality Bipartite Matching (MCBM)

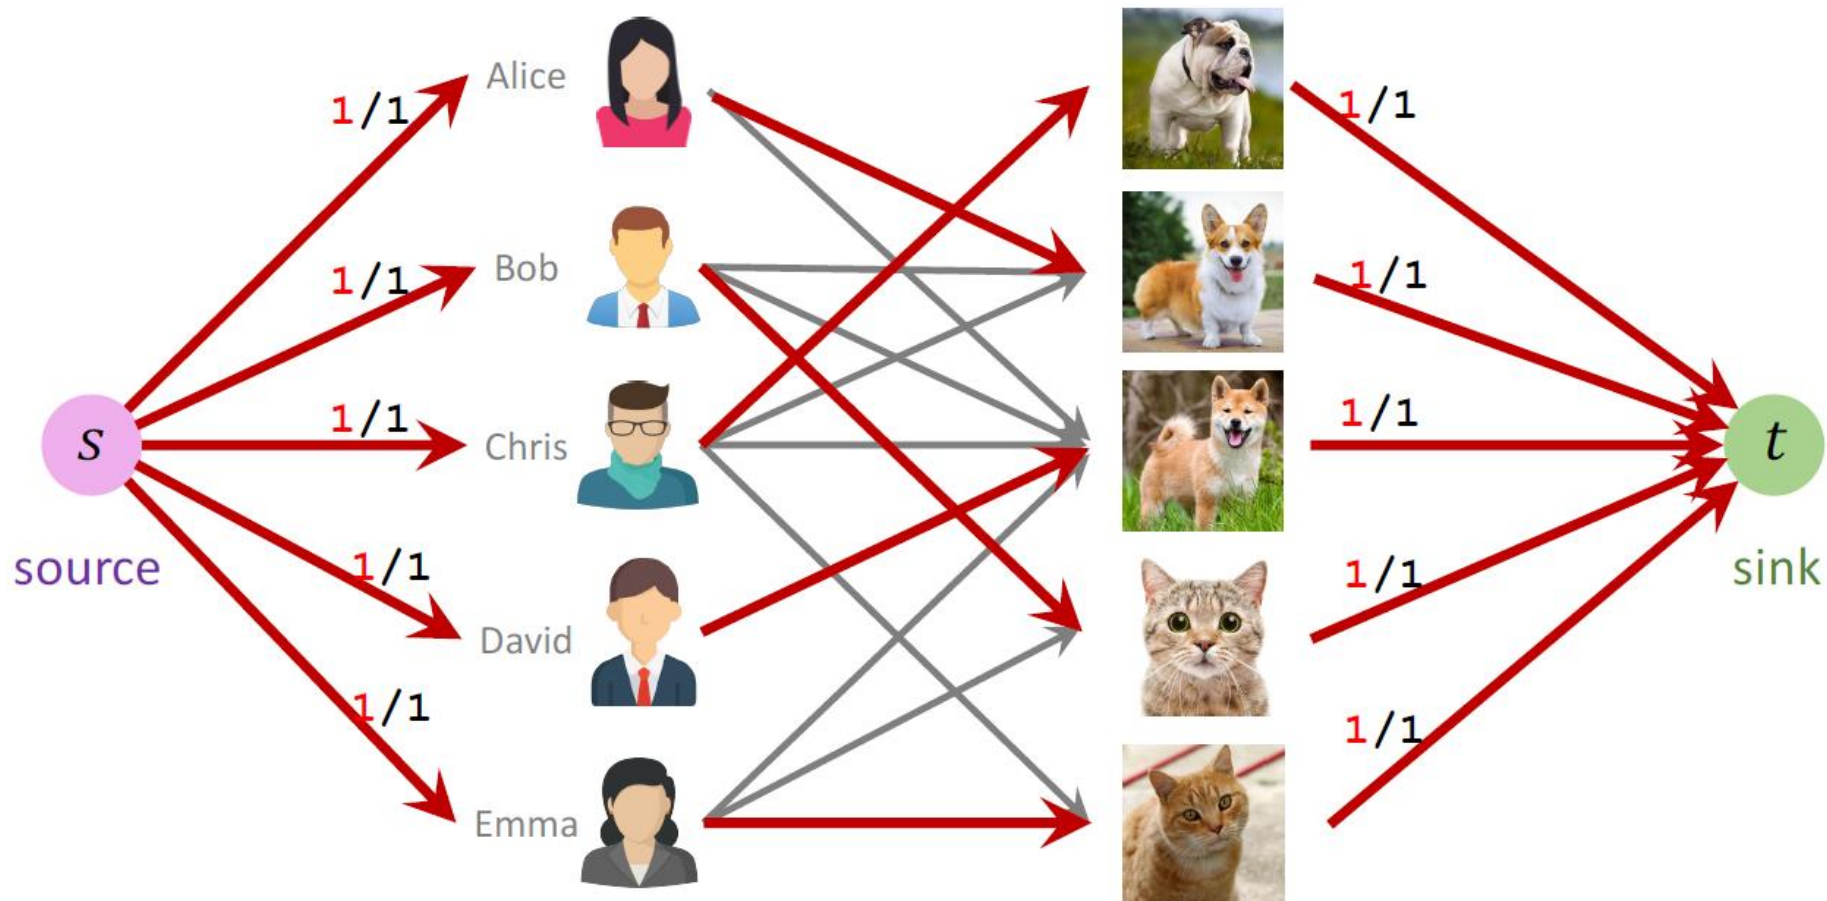
- Bipartite graph: $\mathcal{G} = \mathcal{U}, \mathcal{V}, \mathcal{E}$
- Set \mathcal{U} contains people, Set \mathcal{V} contains pets
- Edges in \mathcal{E} are people's preference
- **Goal:** Maximizing the cardinality of matching, 5

Greedy algorithm can fail!



Maximum-Cardinality Bipartite Matching (MCBM)


- Capacity of max-flow = Cardinality of max-matching






Augmenting path algorithm

- Given a matching
 - An alternating path is a path that begins with an unmatched vertex and whose edges belong alternately to the matching and not to the matching
 - An augmenting path is an alternating path that starts from and ends on (two different) unmatched vertices
 - The path length(total number of edges) of augmenting path must be an odd number
- Total number of unmatched edges is always greater than the number of matched edges by 1
- Switch between unmatched edges and matched edges in the augmenting path




- 2 matchings (1→5, 4→7), so that vertices 1, 4, 5, 7 are matched vertices.
- starting from an unmatched vertex 8
- unmatched edge → matched edge → unmatched edge...



Augmenting path algorithm



3 -> 5 cannot, thus 3 -> 6




4 -> 7 -> 1 -> 5 -> 2 cannot

4 -> 8


```
Graph g_;
int num_vertices_U_, num_vertices_V_;
int total_num_;
std::vector<bool> visited_;
std::map<int, int> matching_; // init as -1

bool FindMatching(int u) {
    // iterating sets V
    for (unsigned int v = num_vertices_U_; v < total_num_; v++) {
        // if there is a edge between vertices from U and V
        bool is_connected =
            std::find(g_.adj_list_[u].cbegin(), g_.adj_list_[u].cend(), v)
            != g_.adj_list_[u].cend();
        // augmenting path - unmatched -> matched -> ... -> matched -> unmatched
        if (false == visited_[v] && is_connected) {
            visited_[v] = true;
            // if vertex in V is not matched, we will match it
            // if it is matched already, we then go back to set U
            // we will try to figure out if the vertex from U can be
            // matched with another vertex in V
            // remember unmatched -> matched -> unmatched -> .....
            if (-1 == matching_[v] || FindMatching(matching_[v])) {
                matching_[v] = u;
                matching_[u] = v;
                return true;
            }
        }
    }
    return false;
}
```

Minimum-Weight Bipartite Matching: Hungarian Algorithm



	v_1	v_2	v_3
u_1	8	25	50
u_2	50	35	75
u_3	22	48	150



The **minimum** sum of weight is $50 + 35 + 22 = 107$

Subtract Row Minima

	v_1	v_2	v_3
u_1	8	25	50
u_2	50	35	75
u_3	22	48	150

	v_1	v_2	v_3
u_1	8	25	50
u_2	50	35	75
u_3	22	48	150

	v_1	v_2	v_3
u_1	8 -8	25 -8	50 -8
u_2	50 -35	35 -35	75 -35
u_3	22 -22	48 -22	150 -22

	v_1	v_2	v_3
u_1	0	17	42
u_2	15	0	40
u_3	0	26	128

Now, the row minima are zeros

Subtract Column Minima

	v_1	v_2	v_3		v_1	v_2	v_3		v_1	v_2	v_3		v_1	v_2	v_3
u_1	0	17	42	u_1	0	17	42	u_1	0 -0	17 -0	42 -40	u_1	0	17	2
u_2	15	0	40	u_2	15	0	40	u_2	15 -0	0 -0	40 -40	u_2	15	0	0
u_3	0	26	128	u_3	0	26	128	u_3	0 -0	26 -0	128 -40	u_3	0	26	88

Now, the col minima are zeros

Iteration 1

Repeat:

- A. Cover all the zeros with a **minimum** number of lines
- B. Decide whether to stop
 - If n lines are required, the algorithm stops
 - If less than n lines are required, then continue with Step C
- C. Create additional zeros
 - Find the smallest element (denote k) that is not covered by a line ($k=2$)
 - Subtract k from all uncovered elements
 - Add k to all the elements that are covered twice

	v_1	v_2	v_3
u_1	0	17	2
u_2	15	0	0
u_3	0	26	88

	v_1	v_2	v_3
u_1	0	17	2
u_2	15	0	0
u_3	0	26	88

	v_1	v_2	v_3
u_1	0	17	2
u_2	15	0	0
u_3	0	26	88

	v_1	v_2	v_3
u_1	0	15	0
u_2	15	0	0
u_3	0	24	86

	v_1	v_2	v_3
u_1	0	15	0
u_2	17	0	0
u_3	0	24	86

Not optimal!

Iteration 2

Repeat:


- A. Cover all the zeros with a **minimum** number of lines
- B. Decide whether to stop
 - If n lines are required, the algorithm stops
 - If less than n lines are required, then continue with Step C.
- C. Create additional zeros
 - Find the smallest element (denote k) that is not covered by a line
 - Subtract k from all uncovered elements
 - Add k to all the elements that are covered twice

	v_1	v_2	v_3
u_1	0	15	0
u_2	17	0	0
u_3	0	24	86

Minimum number of lines: 3, thus, stop

Output the matching

- Choose a matching among the zeros
- Think of the zeros as edges
- Select zeros if they are the only zeros in row/col



Maximum Matching and Time Complexity

- Hungarian Algorithm for Maximum Matching
 - Idea: Max Matching \rightarrow Min Matching
 - Negate the signs of all the weights
 - It is equivalent to the minimum matching
 - Run the Hungarian algorithm
- Hungarian algorithm finds a minimum-weight bipartite matching.
 - It requires $U = |V| = n$
 - Time complexity: $O(n^3)$