

Logistic Regression



What is Logistic Regression?

- Logistic regression is a statistical method for predicting binary classes
 - Among a group of loan applicant, whether a person is good credit or bad?
 - Given an income level, whether a person will buy an iPhone or not
- The outcome or target variable is dichotomous in nature.
- Dichotomous means there are only two possible classes.



More Examples?

- Software project completion in time: Yes/No
- Marketing: Given a price point whether an item be sold
- Finance, banking: Will a stock gain? Should I give loan to the applicant
- Supply chain: Can meet delivery target or not
- Retail: Buyer / non buyer
- Property market: a house is easy to be sold or not



Logistic Regression

- Let's use a simple dataset with 100 houses (rows), 2 input features (columns) and 1 output to introduce and train our very first logistic regression model.
- Predict whether a house is easy to be sold or not based on the Size of the house and the number of Bedrooms.

Size (sq. ft.)	Bedrooms	Easy to sell or not
1600	5	1 (Yes)
1200	4	1 (Yes)
740	2	0 (No)
...
1091	4	0 (No)

Logistic Regression

Linear regression + Sigmoid function


$$Z_i = b + S_i \cdot w1 + B_i \cdot w2$$

S_i is the **S**ize of i^{th} house; B_i is the no. of **B**edrooms of the i^{th} house
 $w1$ and $w2$ are the **w**eights; b is the bias/intercept

$$A_i = \frac{1}{1 + e^{-Z_i}}$$

Cross Entropy Loss Function

$$Loss_i = L_i = -[Y_i \log(A_i) + (1 - Y_i) \log(1 - A_i)]$$



0.5 is the default threshold.


If predicted A_i (i.e., probability) ≥ 0.5 , rounded to 1.

If predicted A_i (i.e., probability) < 0.5 , rounded to 0.

When true class $Y_i = 1$,
 $L_i = -\log(A_i)$

Want L_i close to 0,
 then A_i should be close to 1


that is, the predicted A_i
 should be close to true $Y_i = 1$



When true class $Y_i = 0$,
 $L_i = -\log(1 - A_i)$

Want L_i close to 0,
 Then A_i should be close to 0

that is, the predicted A_i
 should be close to true $Y_i = 0$



$$\text{Mean Loss} = L = \frac{\sum_{i=1}^n L_i}{n}$$

Sometimes, people also use J to denote mean loss.

Gradient Descent

$$\left\{ \begin{array}{l} Z_i = b + S_i \cdot w1 + B_i \cdot w2 \\ A_i = \frac{1}{1 + e^{-Z_i}} \\ L_i = -[Y_i \log(A_i) + (1 - Y_i) \log(1 - A_i)] \\ L = \frac{\sum_{i=1}^n L_i}{n} \end{array} \right.$$

$$\frac{\partial L}{\partial w1} = \frac{\sum_{i=1}^n (\frac{\partial L_i}{\partial w1})}{n} = \frac{\sum_{i=1}^n (\frac{\partial L_i}{\partial A_i} \cdot \frac{\partial A_i}{\partial Z_i} \cdot \frac{\partial Z_i}{\partial w1})}{n} = \frac{\sum_{i=1}^n (A_i - Y_i) \cdot S_i}{n}$$

Similarly,

$$\frac{\partial L}{\partial w2} = \frac{\sum_{i=1}^n (\frac{\partial L_i}{\partial A_i} \cdot \frac{\partial A_i}{\partial Z_i} \cdot \frac{\partial Z_i}{\partial w2})}{n} = \frac{\sum_{i=1}^n (A_i - Y_i) \cdot B_i}{n}$$


Then,

$$w1_{new} = w1_{old} - \frac{\partial L}{\partial w1} \cdot \alpha \quad w2_{new} = w2_{old} - \frac{\partial L}{\partial w2} \cdot \alpha \quad b_{new} = b_{old} - \frac{\partial L}{\partial b} \cdot \alpha$$

$$\frac{\partial L_i}{\partial A_i} = -\frac{Y_i}{A_i} + \frac{1 - Y_i}{1 - A_i}$$

$$\frac{\partial A_i}{\partial Z_i} = A_i(1 - A_i) \quad \frac{\partial Z_i}{\partial w1} = S_i$$

$$\frac{\partial L}{\partial b} = \frac{\sum_{i=1}^n (\frac{\partial L_i}{\partial A_i} \cdot \frac{\partial A_i}{\partial Z_i} \cdot \frac{\partial Z_i}{\partial b})}{n} = \frac{\sum_{i=1}^n (A_i - Y_i) \cdot 1}{n}$$



Logistic Regression Codes

$$\left\{ \begin{array}{l} Z_i = b + S_i \cdot w_1 + B_i \cdot w_2 \\ A_i = \frac{1}{1 + e^{-Z_i}} \\ L_i = -[Y_i \log(A_i) + (1 - Y_i) \log(1 - A_i)] \\ L = \frac{\sum_{i=1}^n L_i}{n} \end{array} \right.$$

```
def calculate_A(S_B, W, b):
    Z = np.dot(W.T, S_B) + b
    A = sigmoid(Z)
    return A.T

def sigmoid(Z):
    A = 1/(1 + np.exp(-Z))
    return A

def loss_function(Y, A):
    L = - np.sum( np.dot(Y.T, np.log(A)) + np.dot((1 - Y).T, np.log(1 - A)) ) / Y.shape[0]
    return L
```

$$\left\{ \begin{array}{l} \frac{\partial L}{\partial w_1} = \frac{\sum_{i=1}^n (A_i - Y_i) \cdot S_i}{n} \\ \frac{\partial L}{\partial w_2} = \frac{\sum_{i=1}^n (A_i - Y_i) \cdot B_i}{n} \\ \frac{\partial L}{\partial b} = \frac{\sum_{i=1}^n (A_i - Y_i) \cdot 1}{n} \\ w_{1_{new}} = w_{1_{old}} - \frac{\partial L}{\partial w_1} \cdot \alpha \\ w_{2_{new}} = w_{2_{old}} - \frac{\partial L}{\partial w_2} \cdot \alpha \\ b_{new} = b_{old} - \frac{\partial L}{\partial b} \cdot \alpha \end{array} \right.$$

```
def update_weights_bias(W, b, S_B, Y, A, learning_rate):
    dW = np.dot(S_B, (A - Y)) / S_B.shape[1] # 2x1 matrix
    db = np.sum(A - Y) / S_B.shape[1]

    W = W - dW * learning_rate
    b = b - db * learning_rate
    return W, b
```

Logistic Regression Codes

```

1 # Logistic regression training
2 def train(X, Y, iters, learning_rate):
3
4     W = np.zeros((X_train.shape[0], 1)) #initialize weights as 2x1 matrix with value 0
5     b = 0
6
7     for i in range(iters):
8         A = calculate_A(X, W, b)
9         L = loss_function(Y, A)
10        W, b = update_weights_bias(W, b, X, Y, A, learning_rate)
11
12        loss.append(L)
13        if i == 10 or i == 1000 or i == 10000 or i >= iters-4:
14            print ("iter={:d} \t W1={:f} \t W2={:f} \t b={:f} \t loss={:f}".format(i, W[0][0], W[1][0], b, L))
15
16    return W, b

```

```

1 loss = []
2 W, b = train(X_train, Y_train, 200000, 0.01)
3
4 print("\n best W: \n", W, "\n and b:\n", b)
5

```

iter=10	W1=0.010789	W2=-0.003872	b=-0.005466	loss=0.691814
iter=1000	W1=0.941400	W2=-0.236599	b=-0.367267	loss=0.585017
iter=10000	W1=5.499376	W2=-0.487563	b=-2.271309	loss=0.288622
iter=199996	W1=19.776076	W2=-0.091501	b=-8.888023	loss=0.092960
iter=199997	W1=19.776110	W2=-0.091501	b=-8.888038	loss=0.092960
iter=199998	W1=19.776143	W2=-0.091501	b=-8.888053	loss=0.092960
iter=199999	W1=19.776177	W2=-0.091502	b=-8.888068	loss=0.092960

best W:
[[19.77617678]
[-0.0915015]]
and b:
-8.888068381139373

```

1 preds_train = predict(X_train, W, b)
2 print(f"Accuracy on train data: {test_accuracy(preds_train, Y_train)}%")
3
4 preds_test = predict(X_test, W, b)
5 print(f"Accuracy on test data: {test_accuracy(preds_test, Y_test)}%")

```

Accuracy on train data: 98.75%
Accuracy on test data: 95.0%

```

1 sample = np.array([[1200,3]])
2 normalized_sample = scaler.transform(sample)
3 normalized_sample = normalized_sample.T
4
5 print ("For a house with size_sqft = {:d} and "
6        "num_bedrooms = {:d}, the prediction is".format(sample[0][0],sample[0][1]),
7        '1 (i.e., Yes, easy to sell)' if predict(normalized_sample, W, b)[0][0] == 1
8        else '0 (i.e., No, not easy to sell)')

```

For a house with size_sqft = 1200 and num_bedrooms = 3, the prediction is 1 (i.e., Yes, easy to sell)



