



Logistic Regression

Multiclass Classification with Logistic Regression


- We have understood that logistic regression, by nature, works for binary classification.
- What if we have more than 2 classes (i.e., multiclass/multinomial), can logistic regression algorithm still be applied? 
- Answer is Yes!
- Question is how?!

Multiclass Classification with Logistic Regression

One-vs-all (one-vs-rest):



Class 1: **Green**
Class 2: **Blue**
Class 3: **Red**




One-vs-Rest (OvR) approach

- Build k independent logistic regression models for k classes
- For example,
 - 1st model for Class 1 **Green** vs Rest
 - 2nd model for Class 2 **Blue** vs Rest
 - 3rd model for Class 3 **Red** vs Rest
- When a previously unseen instance comes in, to make a prediction, we need to run the 3 models and pick the class with the highest probability.
- Say if $P(\text{Green}) > P(\text{Blue})$ and $P(\text{Green}) > P(\text{Red})$, then predicted class is **Green**.

Multiclass Classification with Logistic Regression

- In the One-vs-Rest (OvR) approach, k models need to be built for k classes, which does not sound very efficient. Can we make some improvement, at least a bit?
- Recall that for 2 classes (e.g., classes 0 and 1), we only need to build 1 logistic regression model. Now for k classes, let's try to build $k - 1$ classes.



Sigmoid function plot

For 2 classes (0 and 1)

$$Z = b + w_1x_1 + w_2x_2 + \dots$$

$$A = \frac{1}{1 + e^{-Z}} \quad \frac{A}{1 - A} = e^Z$$

$$\frac{1}{A} = 1 + e^{-Z} \quad \ln\left(\frac{A}{1 - A}\right) = Z$$

$$\frac{1}{A} - 1 = e^{-Z} \quad \ln\left(\frac{A(1)}{A(0)}\right) = Z$$

$$\frac{1 - A}{A} = \frac{1}{e^Z}$$

$A(1)$ is the probability of a sample belongs to class 1
 $A(0)$ is the probability of a sample belongs to class 0

Now let's say 3 classes (0, 1 and 2) Important constraint:

1st model: class 0 vs 1

$$\ln\left(\frac{A(1)}{A(0)}\right) = Z^{(01)} \\ = b^{(01)} + w_1^{(01)}x_1 + w_2^{(01)}x_2 + \dots$$

$$A(1) = A(0)e^{Z^{(01)}}$$

2nd model: class 0 vs 2

$$\ln\left(\frac{A(2)}{A(0)}\right) = Z^{(02)} \\ = b^{(02)} + w_1^{(02)}x_1 + w_2^{(02)}x_2 + \dots$$

$$A(2) = A(0)e^{Z^{(02)}}$$

$$A(0) + A(1) + A(2) = 1$$

$$A(0) + A(0)e^{Z^{(01)}} + A(0)e^{Z^{(02)}} = 1$$

So,

$$A(0) = \frac{1}{1 + e^{Z^{(01)}} + e^{Z^{(02)}}}$$

$$A(1) = \frac{e^{Z^{(01)}}}{1 + e^{Z^{(01)}} + e^{Z^{(02)}}}$$

$$A(2) = \frac{e^{Z^{(02)}}}{1 + e^{Z^{(01)}} + e^{Z^{(02)}}}$$

0.5 is the default threshold.

If predicted A_i (i.e., probability) ≥ 0.5 , rounded to 1.

If predicted A_i (i.e., probability) < 0.5 , rounded to 0.

Regularization

- $\min_{w,b} J(w,b)$, where $J = L$, another common symbol to denote loss/error in ML
- Parameters: $w \in \mathbb{R}^{n_x}$, $b \in \mathbb{R}$, λ = regularization parameter, n_x = no. of features

- $J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \|w\|_2^2$
 w is a vector of all weights

$\hat{y}^{(i)}$ is the predicted value of the i^{th} sample
 $y^{(i)}$ is the true value of the i^{th} sample
 L is the Cross Entropy Loss function we used previously

- L_2 regularization: $\|w\|_2^2 = \sum_{j=1}^{n_x} w_j^2 = w^T w$

- L_1 regularization: $\|w\|_1 = \sum_{j=1}^{n_x} |w_j|$

$$\sum_{j=1}^{n_x} w_j^2 = w^T w = [w_1 \ w_2 \ w_3] \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix}$$

$$= w_1^2 + w_2^2 + w_3^2$$

$$\sum_{j=1}^{n_x} |w_j| = |w_1| + |w_2| + |w_3|$$

Regularization

Loss function
with L2
regularization

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \|w\|_2^2$$

Find derivatives of loss
w.r.t weights, similar to
what we did previously in
the selling house
example.

Only differences lies in
the introduction of L2
regularization term here

$$\frac{\partial J}{\partial w} = dw = (\text{from backprop}) + \frac{\lambda}{m} w_{old}$$

$$w_{new} = w_{old} - \alpha dw = w_{old} - \alpha [(\text{from backprop}) + \frac{\lambda}{m} w_{old}]$$

$$w_{new} = w_{old} - \alpha \frac{\lambda}{m} w_{old} - \alpha (\text{from backprop})$$

$$w_{new} = (1 - \alpha \frac{\lambda}{m}) w_{old} - \alpha (\text{from backprop})$$

<1, so called "Weight decay"

Gradient
descent

Regularization

- Essentially what regularization does is adding extra penalty to complicated model with higher values of weights.
- That's why regularization is an efficient technique to prevent overfitting and also an important hyperparameter in ML.
- Additional Reading
 - The *penalty* hyperparameter in [sklearn logistic regression](#)
 - For linear regression, check [Lasso](#) for L1 regularization and [Ridge](#) for L2 regularization

